

[537] Virtual Memory

Tyler Harter
9/15/14

Overview

Review Scheduling

Address Spaces (Chapter 13)

Address Translation (Chapter 15)

Segmentation (Chapter 16)

Review: Schedulers

Scheduling Basics

Workloads:

arrival_time

run_time

Schedulers:

FIFO

SJF

STCF

RR

Metrics:

turnaround_time

response_time

Scheduling Basics

Workloads:

arrival_time
run_time

Schedulers:

FIFO
SJF
STCF
RR

Metrics:

turnaround_time
response_time


**Project grading
will be based on
turnaround time!**

Workload

JOB	arrival	run
A	0	40
B	0	20
C	5	10

Timelines

ABCABCABABAAAA



B C A



Schedulers:


FIFO

SJF


STCF

RR

BC B A



A B C



Workload

JOB	arrival	run
A	0	40
B	0	20
C	5	10

Schedulers:

FIFO


SJF

STCF

RR


Timelines

ABCABCABABAAA




RR

B C A




SJF

BC B A




STCF

A B C



FIFO

Lottery Scheduler




```
winner = random(402)
```



winner = 102

Job A (1)	Job B (1)	Job C (100)	Job D (200)	Job E (100)
--------------	--------------	----------------	----------------	----------------

← **402** total tickets →

winner = 102

is 102 < 1 ?



Job A (1)	Job B (1)	Job C (100)	Job D (200)	Job E (100)
--------------	--------------	----------------	----------------	----------------

← 402 total tickets →

winner = 101

is 101 < 1 ?



Job A (1)	Job B (1)	Job C (100)	Job D (200)	Job E (100)
--------------	--------------	----------------	----------------	----------------

← 402 total tickets →

winner = 100

is 100 < 100 ?



Job A (1)	Job B (1)	Job C (100)	Job D (200)	Job E (100)
--------------	--------------	----------------	----------------	----------------

← 402 total tickets →

winner = 0

is 0 < 200 ?



Job A (1)	Job B (1)	Job C (100)	Job D (200)	Job E (100)
--------------	--------------	----------------	----------------	----------------

← 402 total tickets →

Run D!

is $0 < 200$?



← **402** total tickets →

Address Spaces

More Virtualization

Virtual CPU: *illusion* of private CPU registers

- 2 lectures

Virtual RAM: *illusion* of private memory

- 5 lectures

The 1st “Easy Piece” in OSTEP is virtual CPU+RAM

The Abstraction

A process has a set of addresses that **map** to bytes

This set is called on **address space**

How can we provide a **private** address space?

Extend LDE (limited direct execution)

Review: what stuff is in an address space?


Match that Segment!

```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof int);  
}
```

x	code
main	data
y	heap
z	stack


Match that Segment!

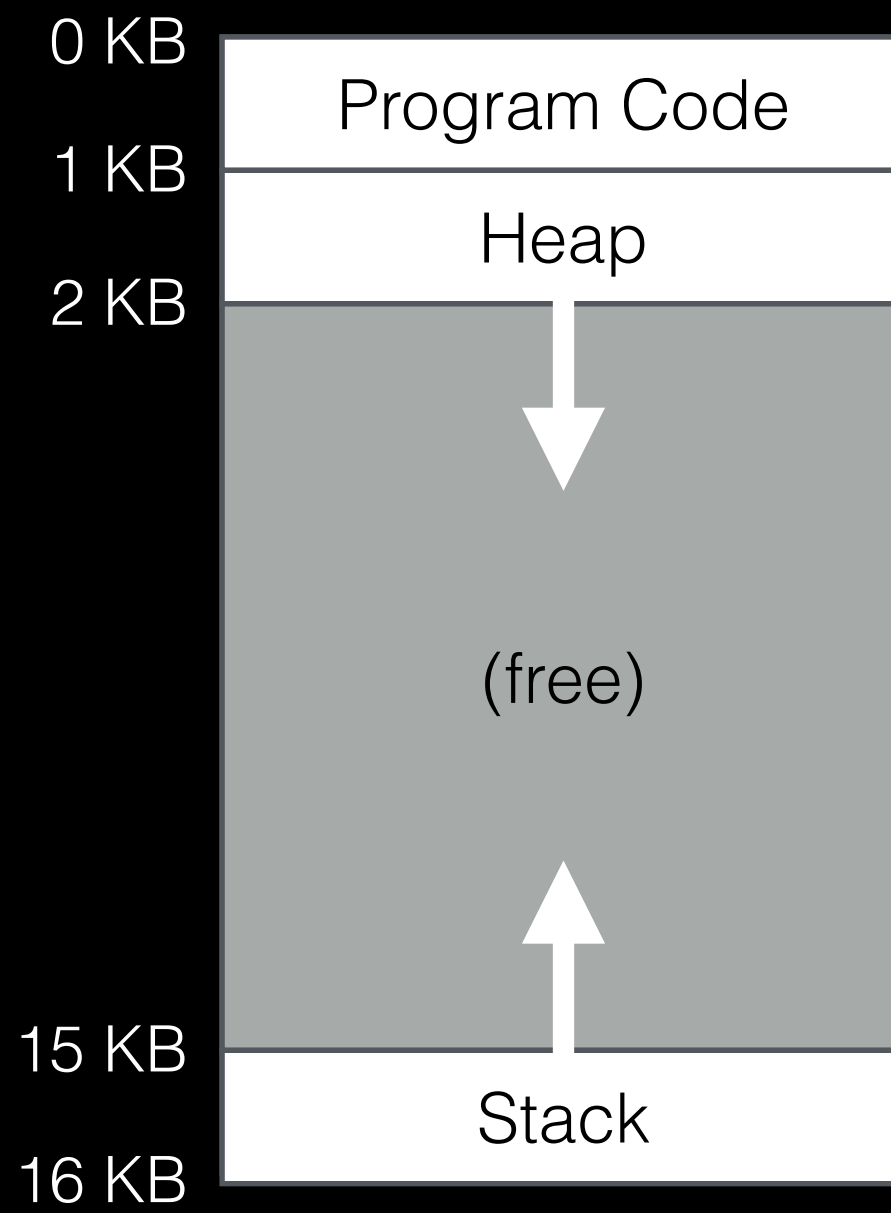
```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof int);  
}
```



Match that Segment!

```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof int);  
}
```





demo0.c output


where is code?	0x100f2ddd0	(4 GB)
where is data?	0x100f2e020	(4 GB)
where is heap?	0x7ff659403930	(131033 GB)
where is stack?	0x7fff5ecd2a1c	(131069 GB)

demo1.c disassemble

Memory Accesses

```
#include <stdio.h>
#include <stdlib.h>
```

```
int main(int argc, char *argv[])
{
    int x;
    x = x + 3;
}
```




```
_main:
0000000000000000 pushq %rbp
0000000000000001 movq %rsp, %rbp
0000000000000004 movl $0x0, %eax
0000000000000009 movl %edi, 0xffffffffc(%rbp)
000000000000000c movq %rsi, 0xffffffff0(%rbp)
0000000000000010 movl 0xffffffffc(%rbp), %edi
0000000000000013 addl $0x3, %edi
0000000000000019 movl %edi, 0xffffffffc(%rbp)
000000000000001c popq %rbp
000000000000001d ret
```

otool -tv demo1.o
(or objdump on Linux)

Memory Accesses

```
#include <stdio.h>
#include <stdlib.h>
```

```
int main(int argc, char *argv[])
{
    int x;
    x = x + 3;
}
```



```
_main:
0000000000000000 pushq %rbp
0000000000000001 movq %rsp, %rbp
0000000000000004 movl $0x0, %eax
0000000000000009 movl %edi, 0xffffffff(%rbp)
000000000000000c movq %rsi, 0xffffffff0(%rbp)
0000000000000010 movl 0xffffffff(%rbp), %edi
0000000000000013 addl $0x3, %edi
0000000000000019 movl %edi, 0xffffffff(%rbp)
000000000000001c popq %rbp
000000000000001d ret
```

otool -tv demo1.o
(or objdump on Linux)

Memory Accesses

%rip = 0x10
%rbp = 0x200

Memory Accesses:


➔ 0x10: movl 0x8(%rbp), %edi
0x13: addl \$0x3, %edi
0x19: movl %edi, 0x8(%rbp)

Memory Accesses

%rip = 0x10
%rbp = 0x200

Memory Accesses:

Fetch instruction at addr 0x10




```
0x10: movl 0x8(%rbp), %edi  
0x13: addl $0x3, %edi  
0x19: movl %edi, 0x8(%rbp)
```

Memory Accesses

%rip = 0x10

%rbp = 0x200



```
0x10: movl 0x8(%rbp), %edi  
0x13: addl $0x3, %edi  
0x19: movl %edi, 0x8(%rbp)
```

Memory Accesses:


Fetch instruction at addr 0x10

Exec, load from addr 0x208

Memory Accesses

%rip = 0x13

%rbp = 0x200



```
0x10: movl 0x8(%rbp), %edi  
0x13: addl $0x3, %edi  
0x19: movl %edi, 0x8(%rbp)
```

Memory Accesses:


Fetch instruction at addr 0x10

Exec, load from addr 0x208

Memory Accesses

%rip = 0x13

%rbp = 0x200



```
0x10: movl 0x8(%rbp), %edi  
0x13: addl $0x3, %edi  
0x19: movl %edi, 0x8(%rbp)
```

Memory Accesses:

Fetch instruction at addr 0x10


Exec, load from addr 0x208

Fetch instruction at addr 0x13

Memory Accesses

%rip = 0x13

%rbp = 0x200



```
0x10: movl 0x8(%rbp), %edi  
0x13: addl $0x3, %edi  
0x19: movl %edi, 0x8(%rbp)
```

Memory Accesses:

Fetch instruction at addr 0x10

Exec, load from addr 0x208

Fetch instruction at addr 0x13


Exec, no load

Memory Accesses

%rip = 0x19

%rbp = 0x200

0x10: movl 0x8(%rbp), %edi
0x13: addl \$0x3, %edi
0x19: movl %edi, 0x8(%rbp)



Memory Accesses:

Fetch instruction at addr 0x10

Exec, load from addr 0x208

Fetch instruction at addr 0x13


Exec, no load

Memory Accesses

%rip = 0x19

%rbp = 0x200

0x10: movl 0x8(%rbp), %edi
0x13: addl \$0x3, %edi
0x19: movl %edi, 0x8(%rbp)



Memory Accesses:

Fetch instruction at addr 0x10

Exec, load from addr 0x208

Fetch instruction at addr 0x13

Exec, no load


Fetch instruction at addr 0x19

Memory Accesses

%rip = 0x19

%rbp = 0x200

0x10: movl 0x8(%rbp), %edi
0x13: addl \$0x3, %edi
0x19: movl %edi, 0x8(%rbp)



Memory Accesses:

Fetch instruction at addr 0x10

Exec, load from addr 0x208

Fetch instruction at addr 0x13

Exec, no load

Fetch instruction at addr 0x19

Exec, store to addr 0x208

Problem: How to Run Multiple Processes?

Addresses are “hardcoded” into process binaries.
How to avoid collisions?

Approaches (covered today):

- Time Sharing

- Static Relocation

- Base


- Base+Bounds

- Segmentation

Time Sharing


We give the illusion of many virtual CPUs by saving **CPU registers** to **memory** when a process isn't running



We give the illusion of many virtual memories by saving **memory** to **disk** when a process isn't running

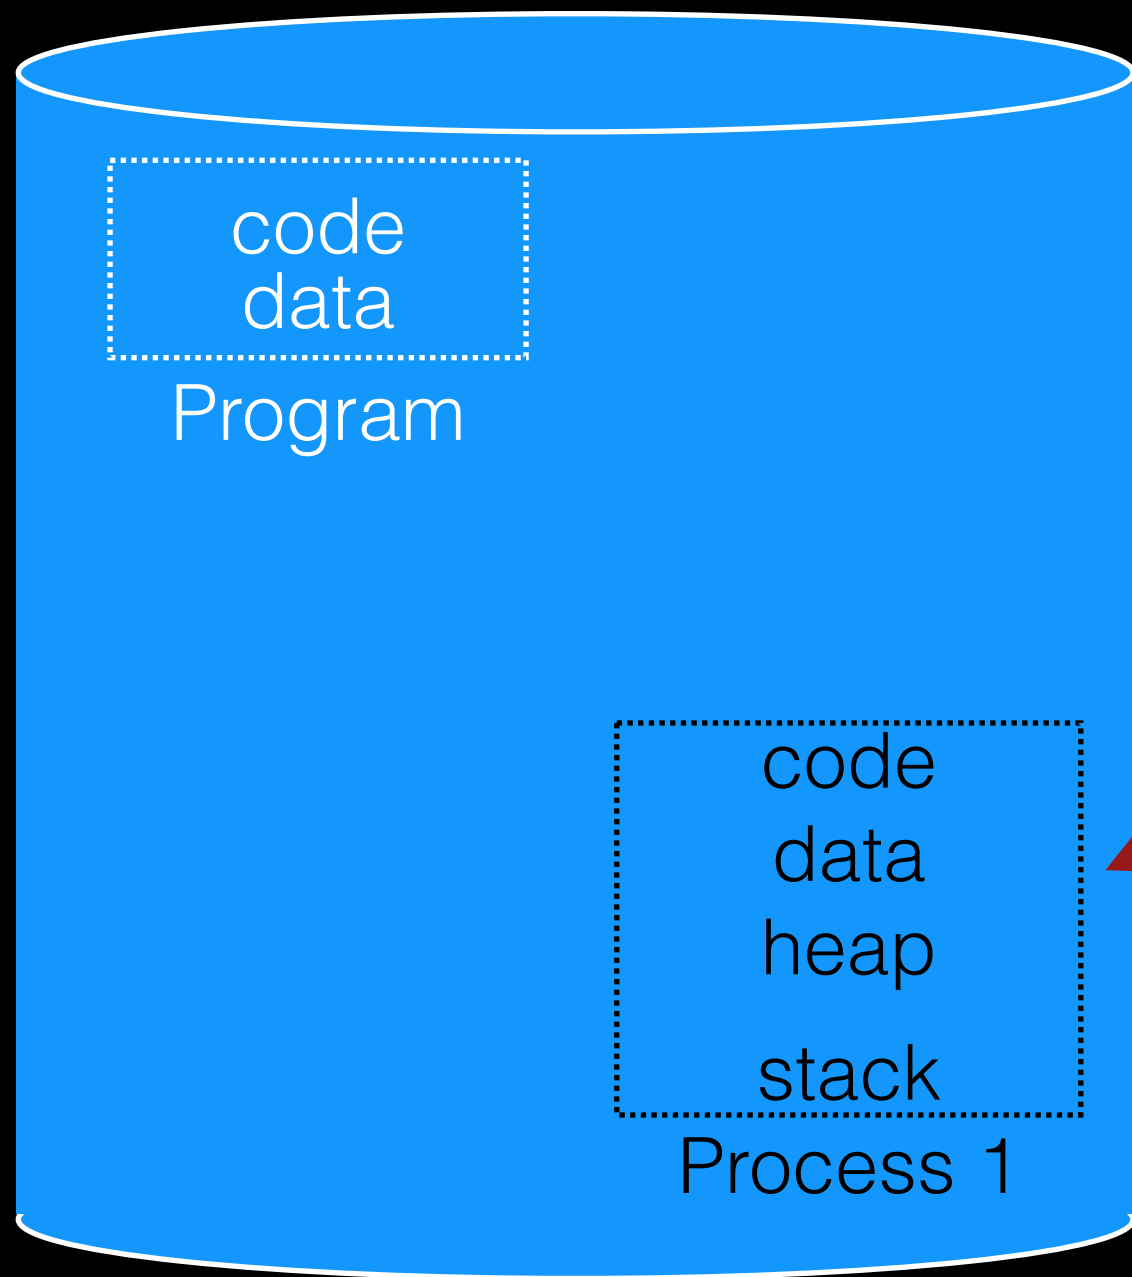


Memory



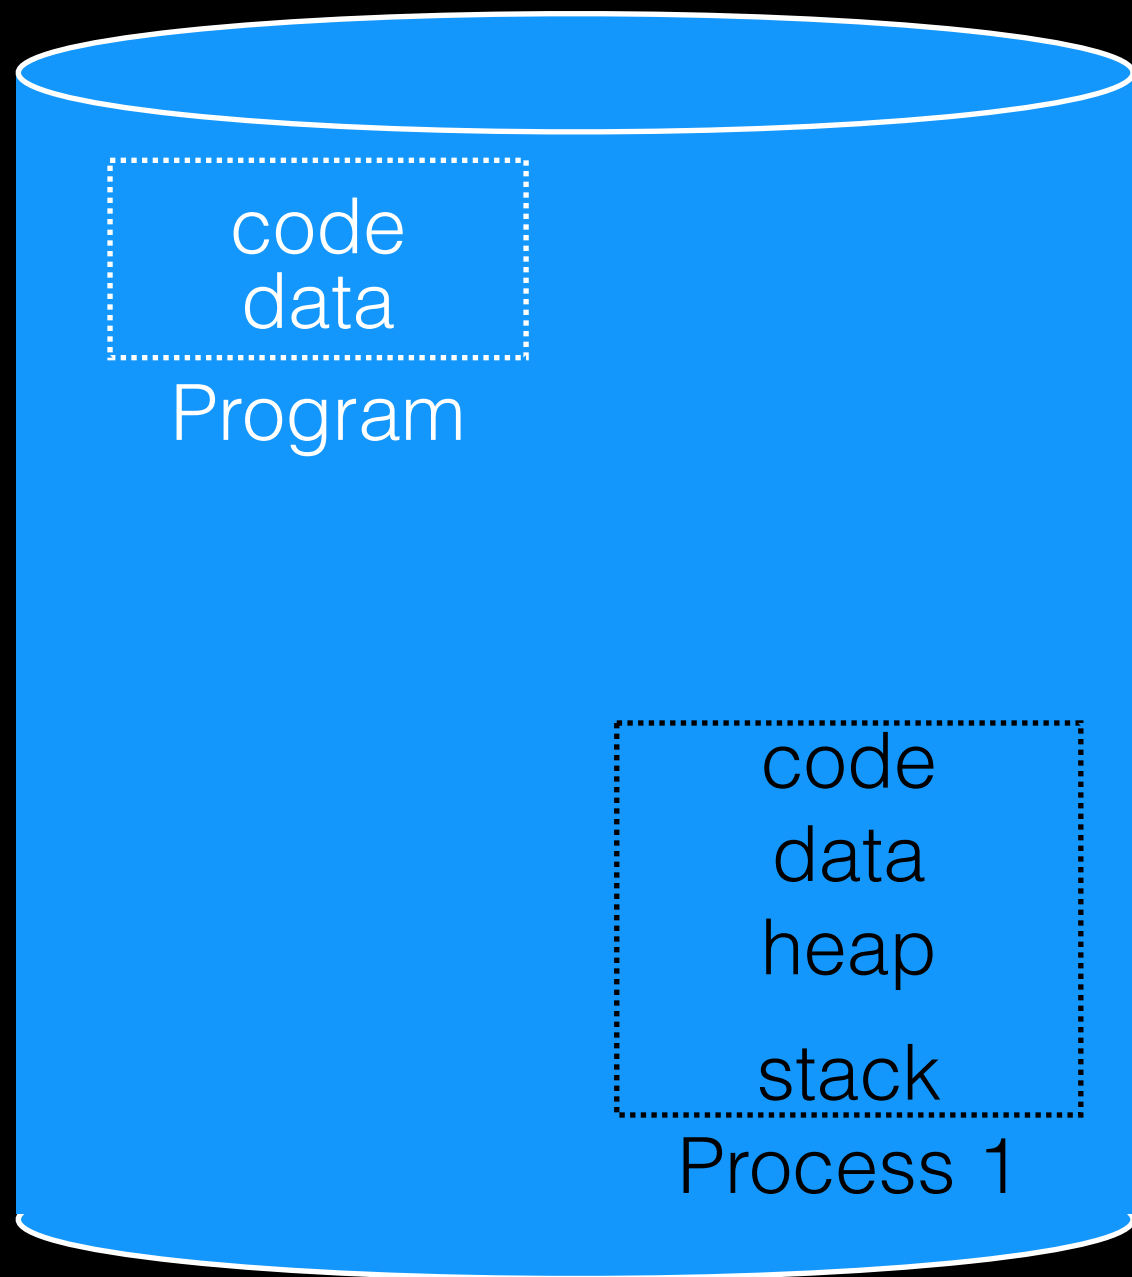






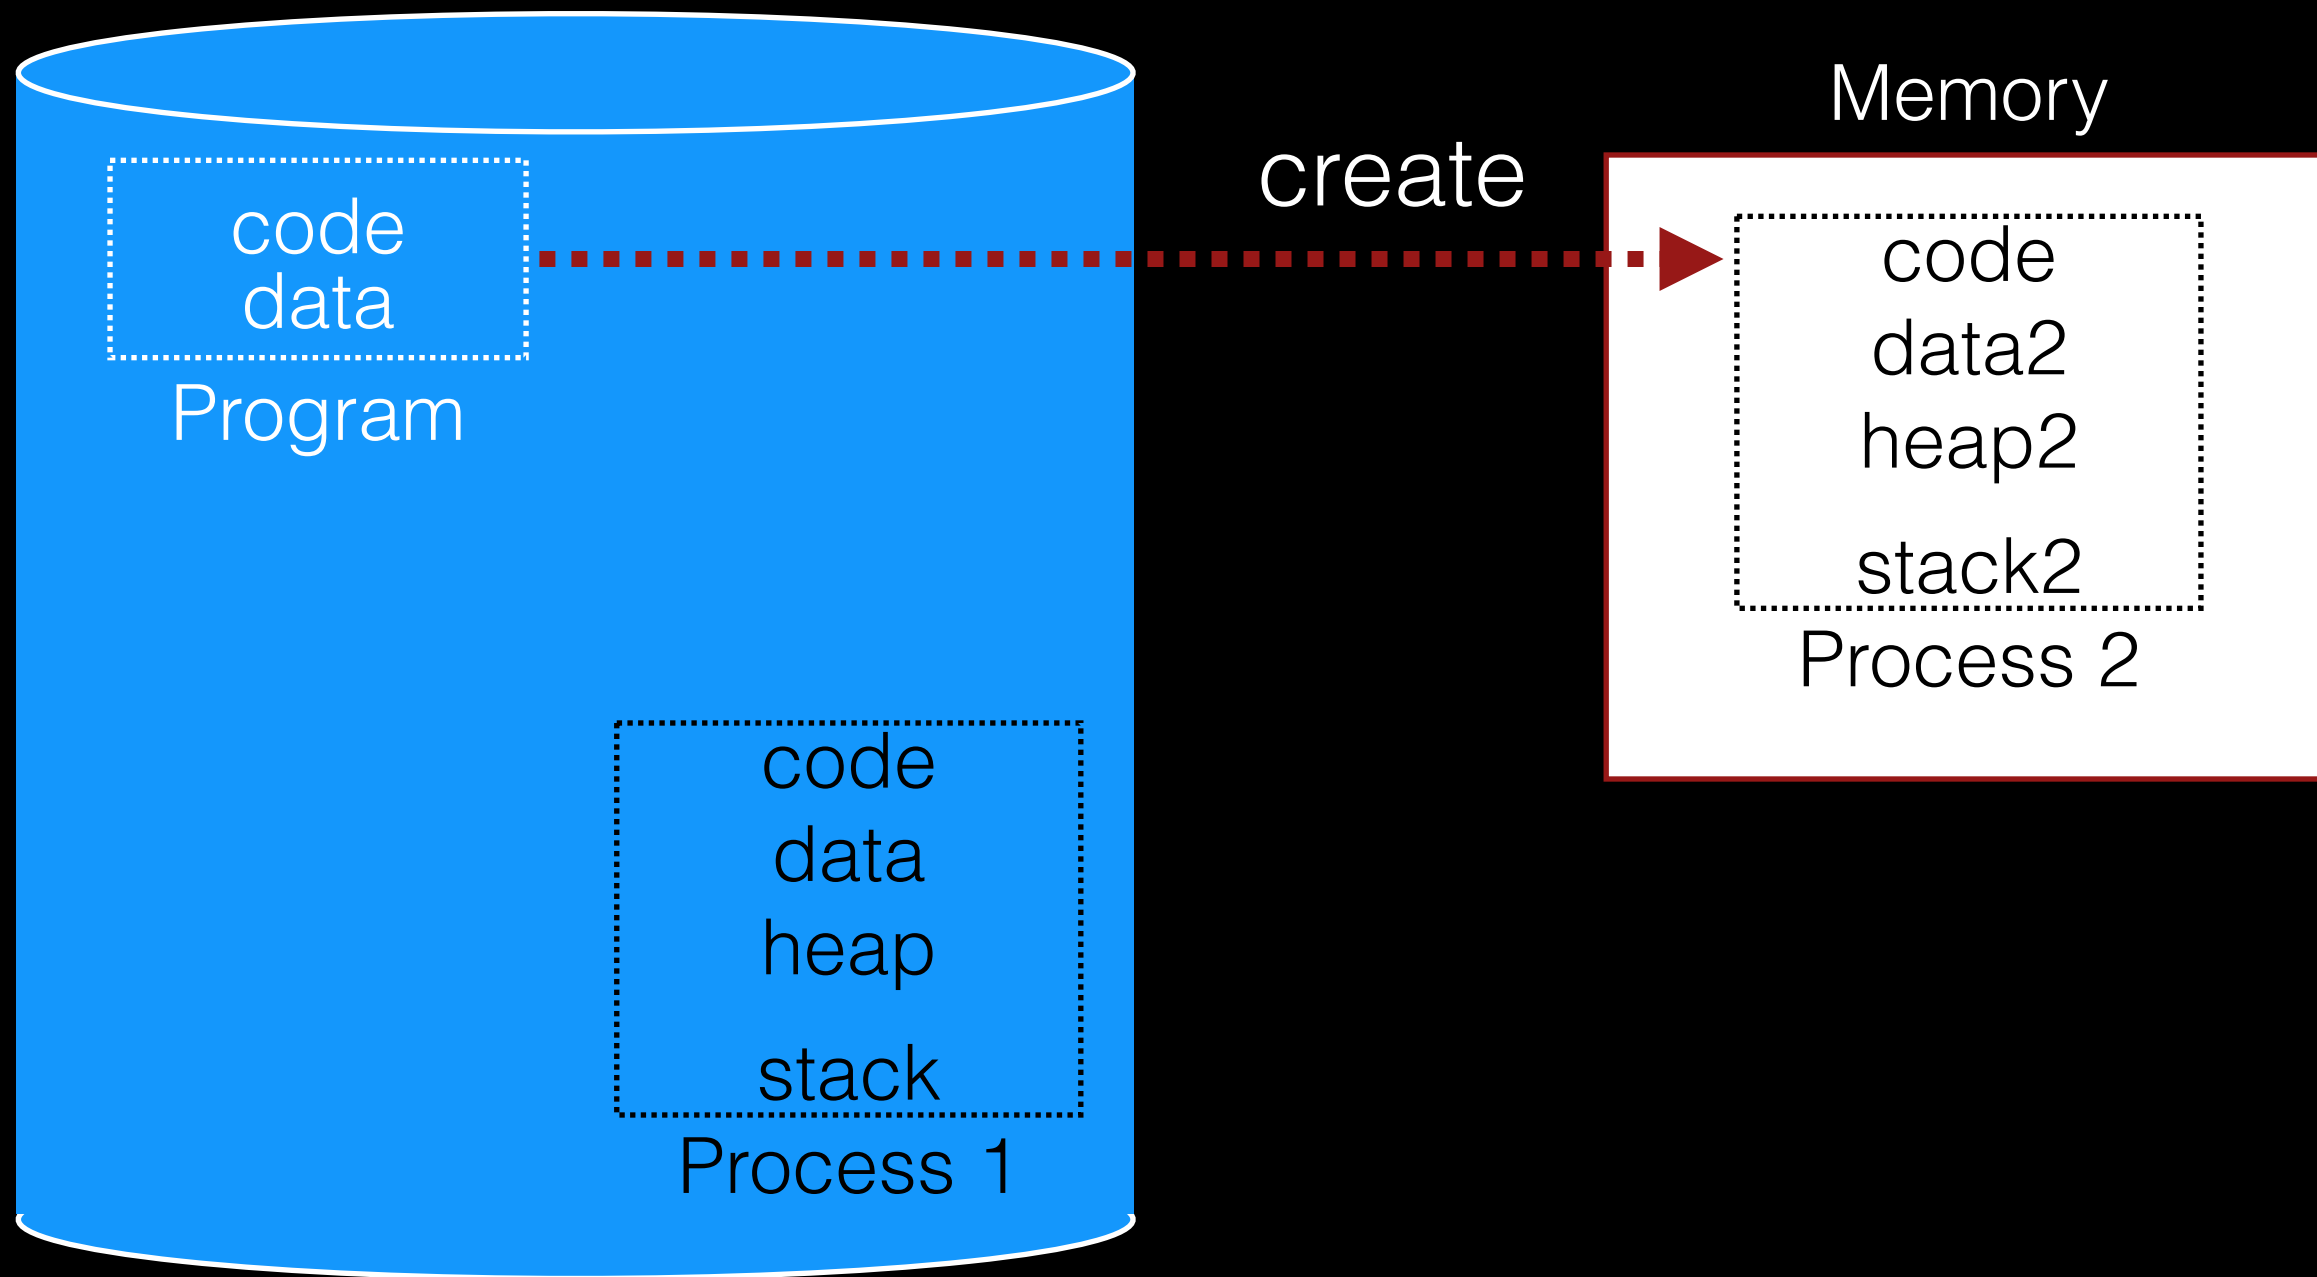
Memory

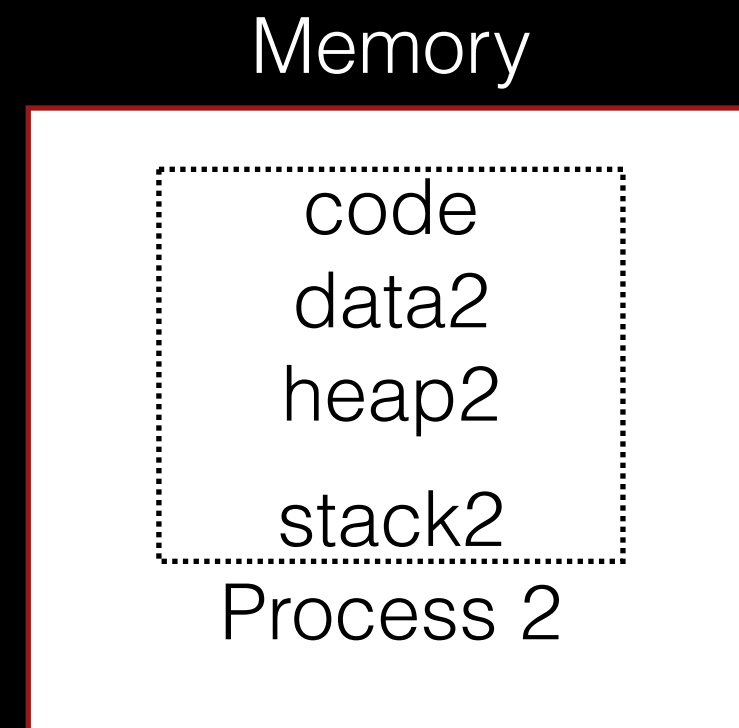
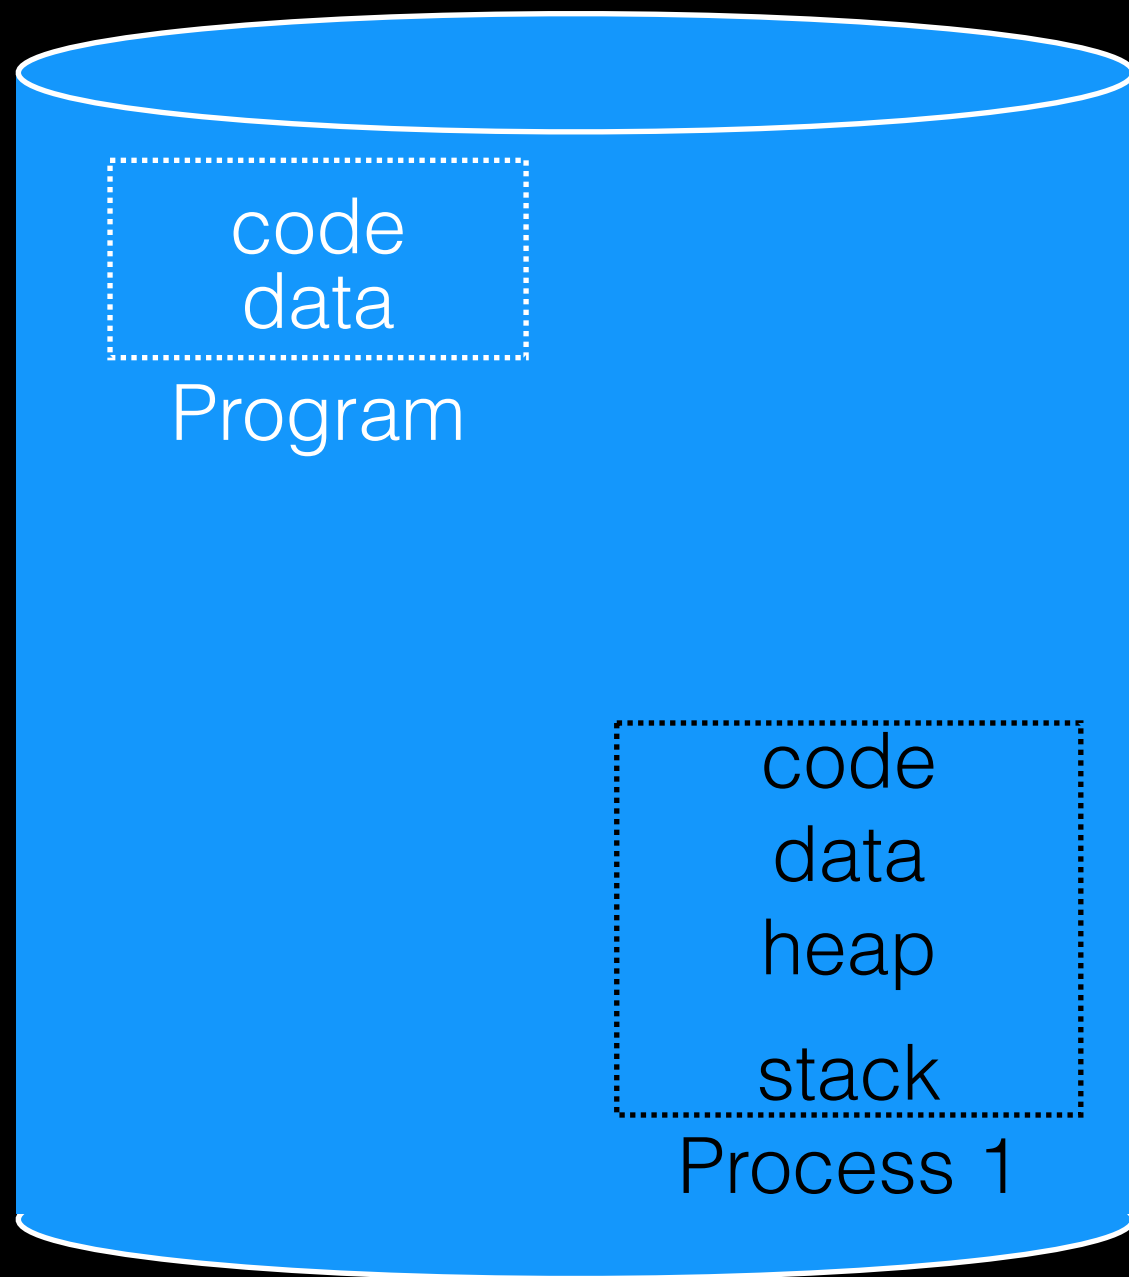


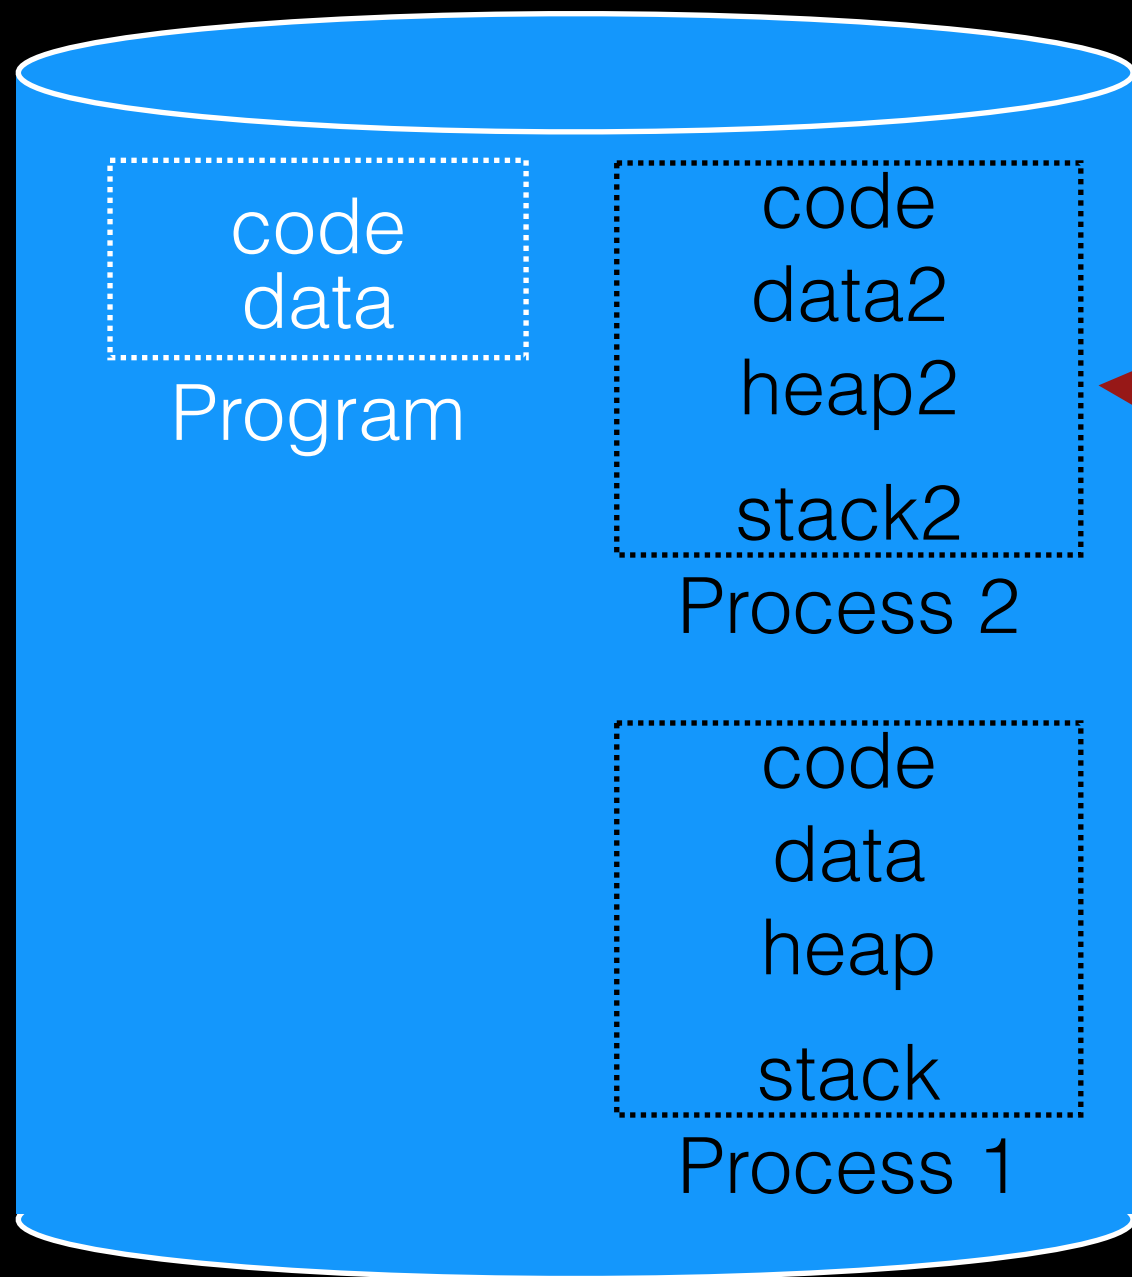


Memory



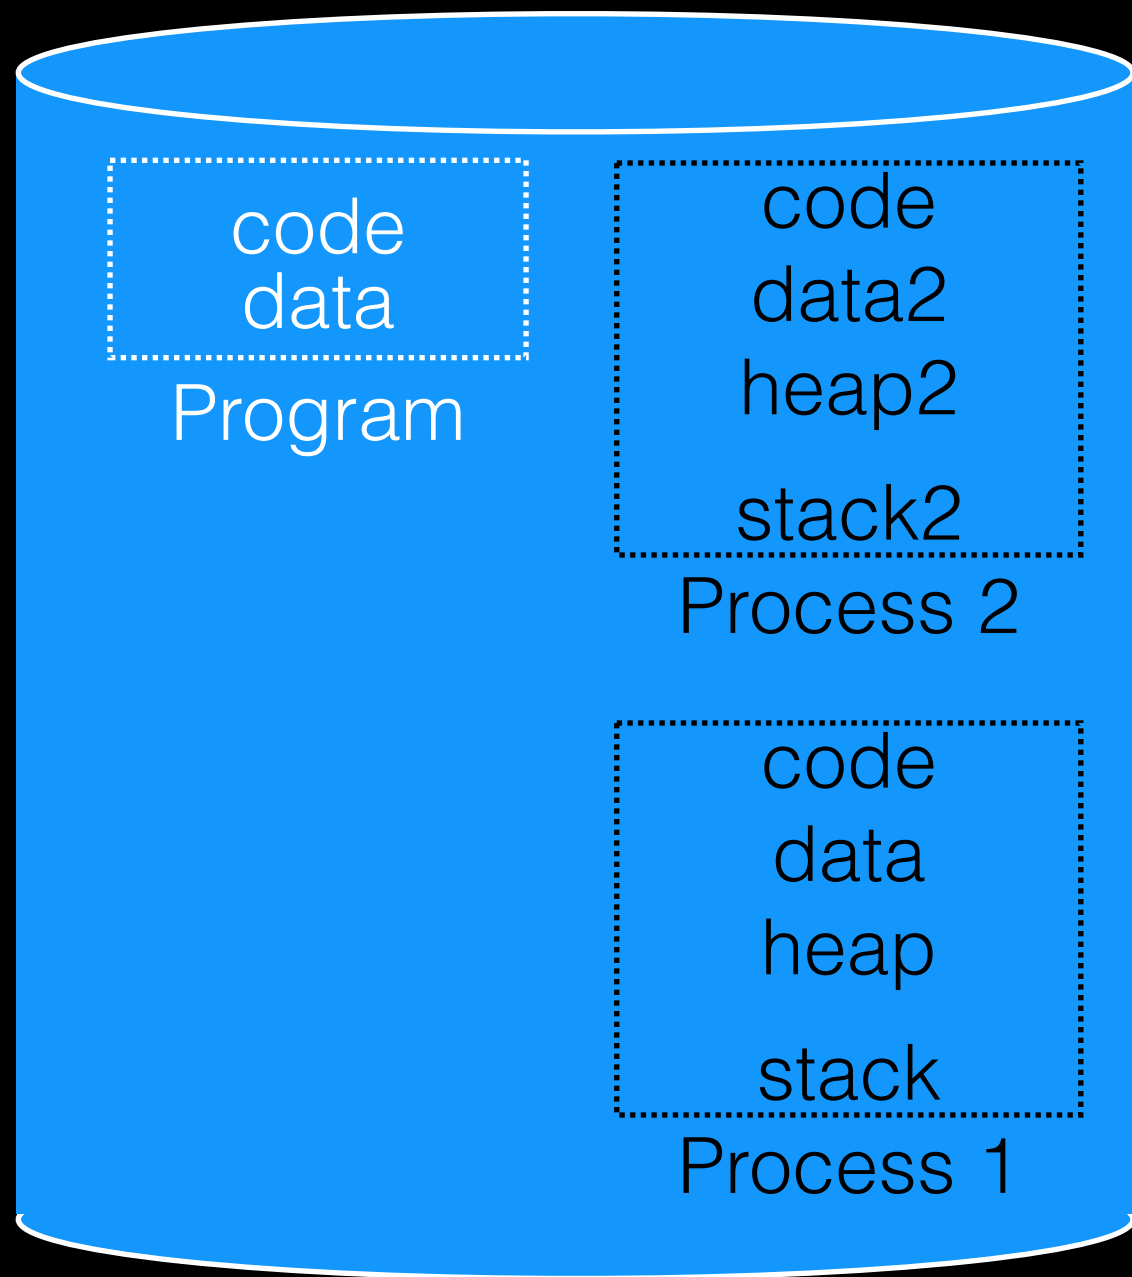







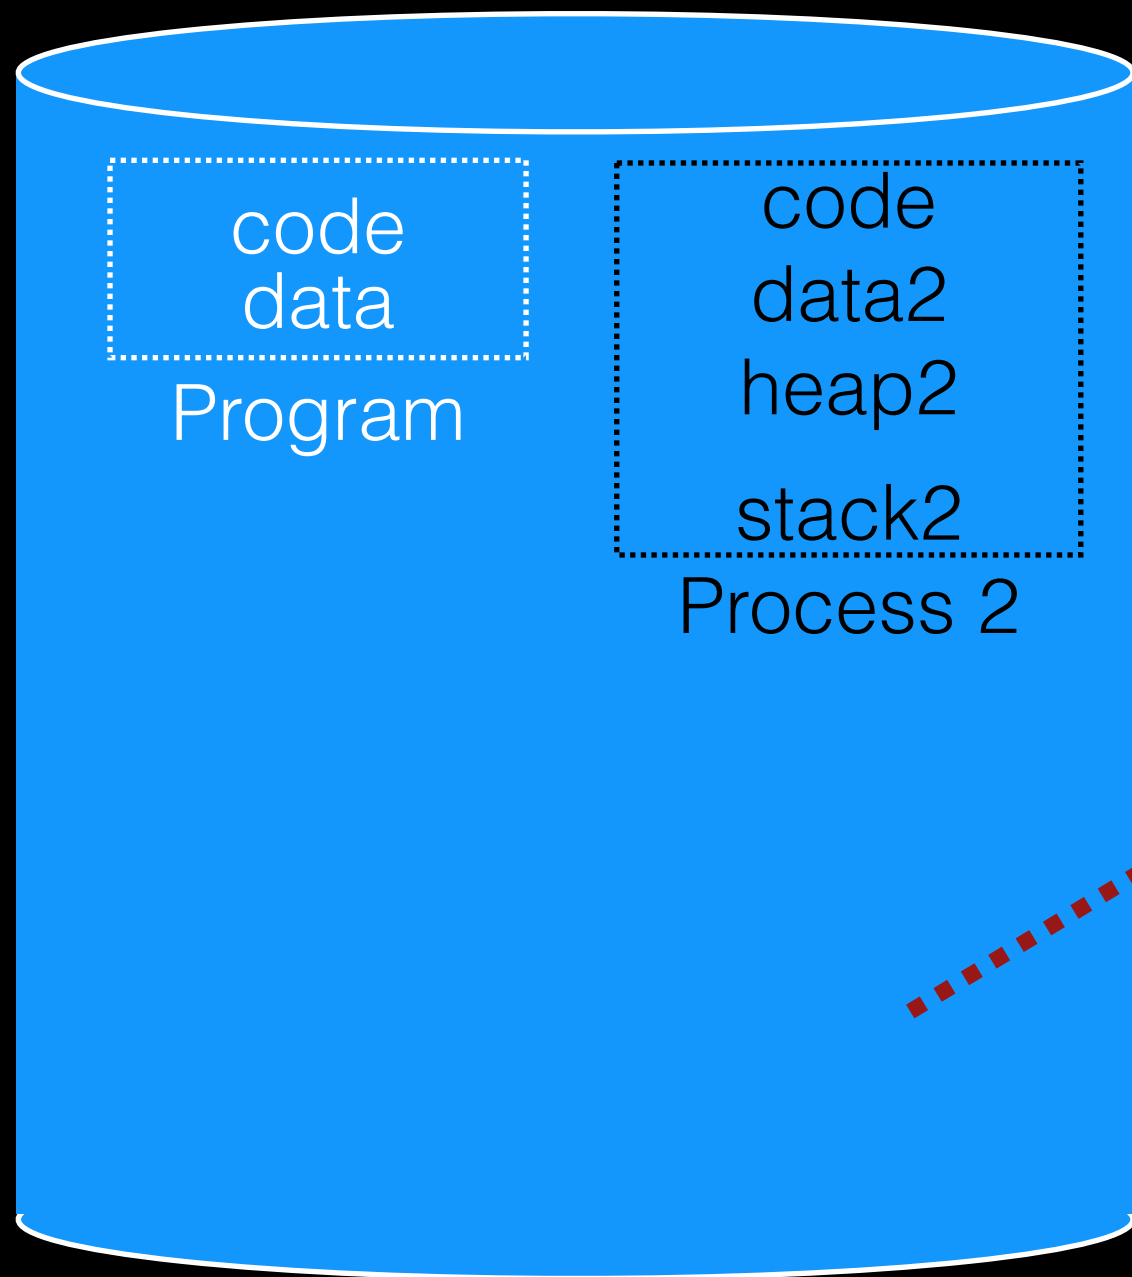
Memory


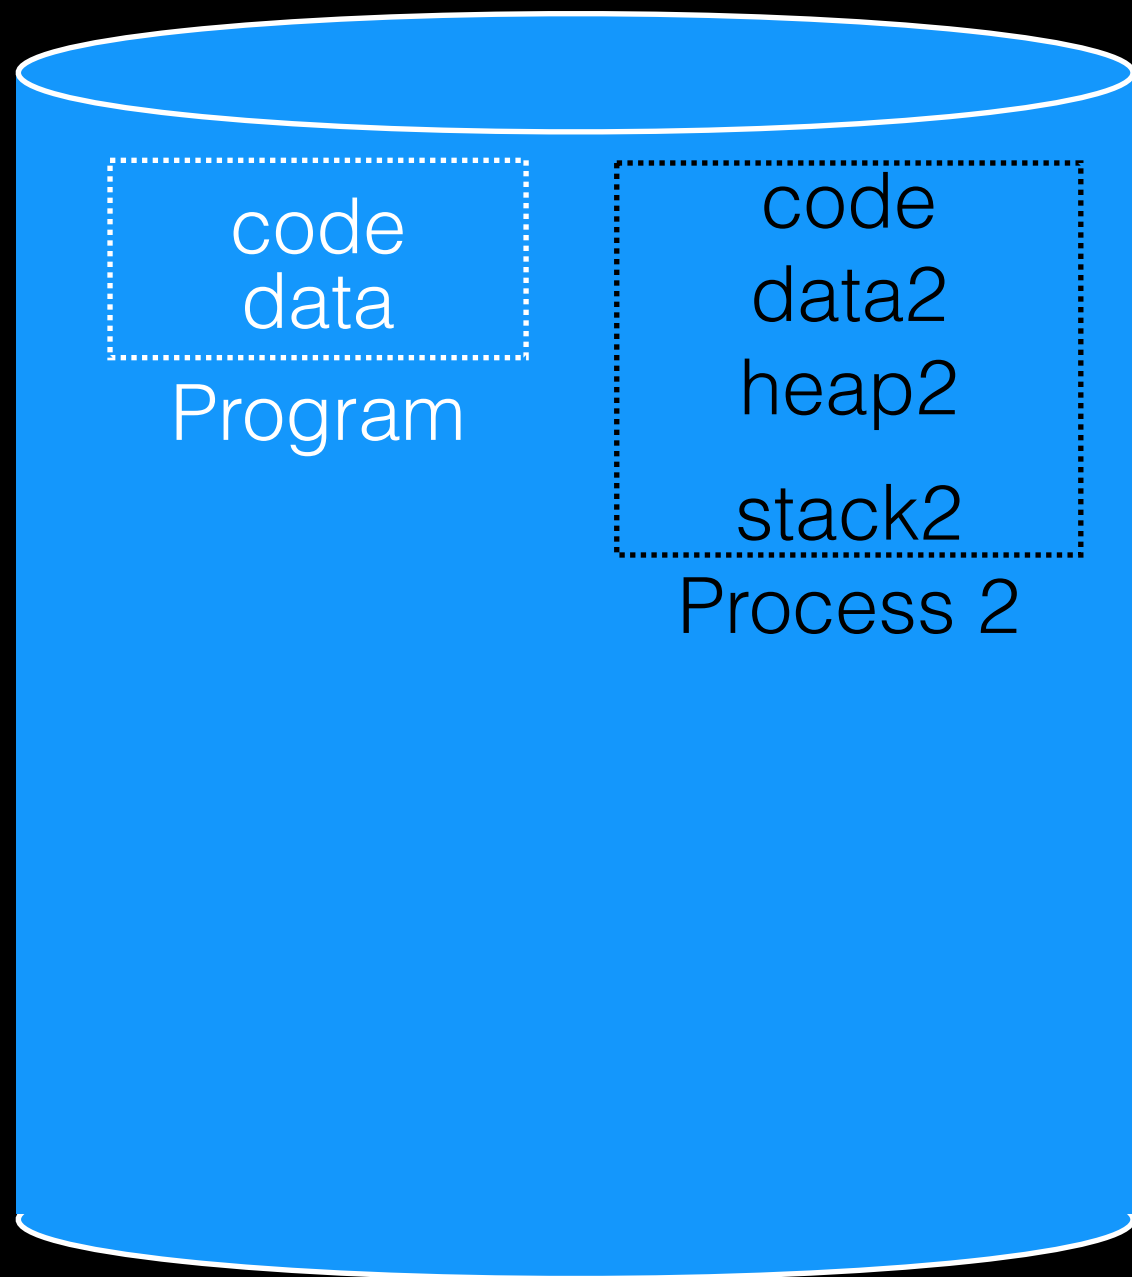




Memory







Time Sharing

Problems?

What schedulers would time sharing work well with?

Alternative: space sharing

Problem: How to Run Multiple Processes?

Approaches (covered today):

Time Sharing

Static Relocation

Base

Base+Bounds

Segmentation

Static Relocation

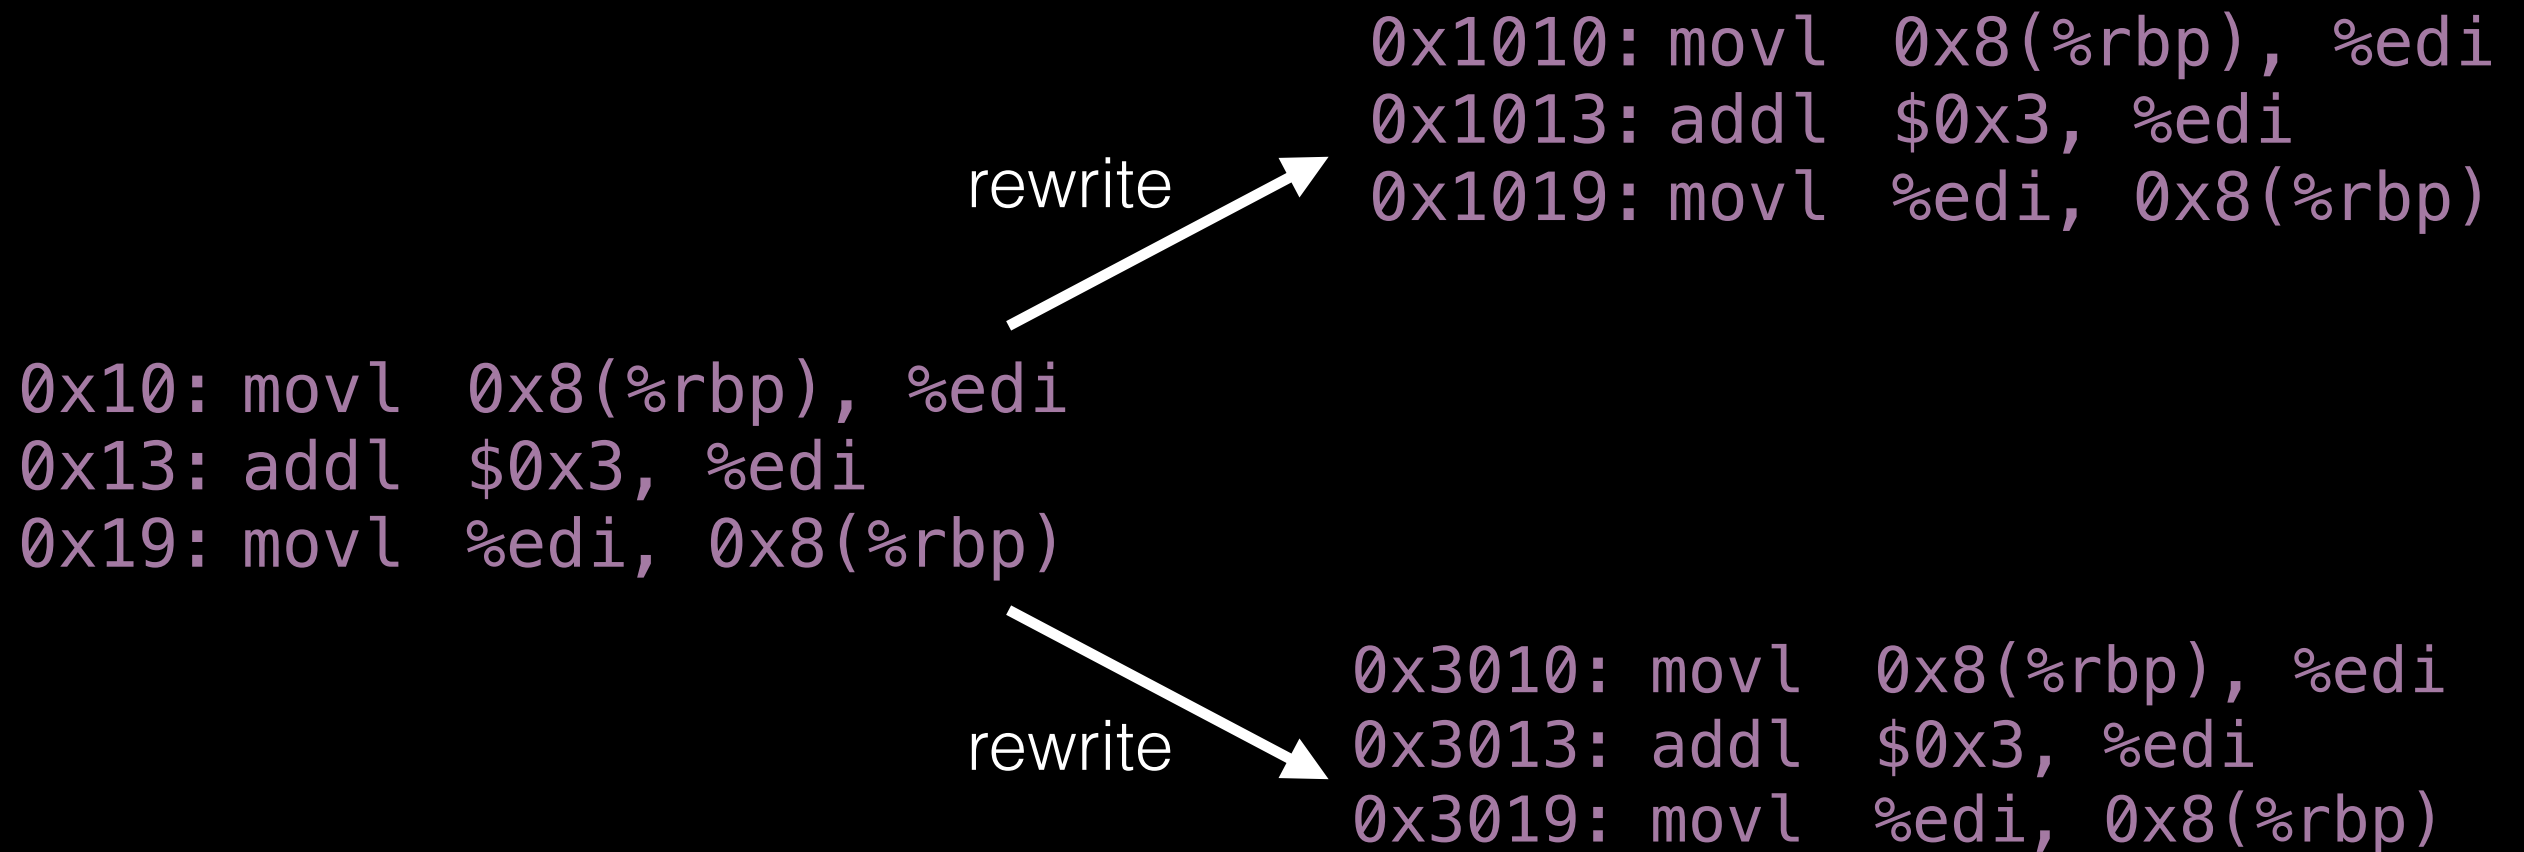
Idea: rewrite each program before loading it as a process

Each rewrite uses different addresses and pointers

Change jumps, loads, etc.

Can any addresses be unchanged?

Rewrite for Each New Process



process 1



4 KB

(free)

Program Code

Heap

(free)

8 KB

stack

(free)

```
0x1010: movl 0x8(%rbp), %edi
0x1013: addl $0x3, %edi
0x1019: movl %edi, 0x8(%rbp)
```

process 2



12 KB

Program Code

Heap

(free)

16 KB

stack

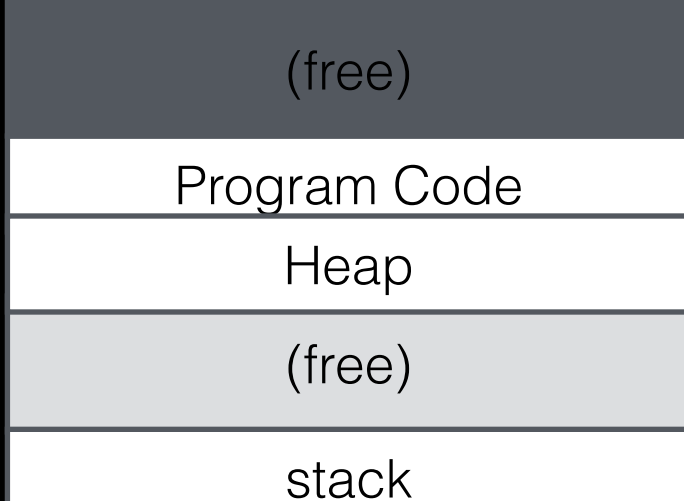
(free)

```
0x3010: movl 0x8(%rbp), %edi
0x3013: addl $0x3, %edi
0x3019: movl %edi, 0x8(%rbp)
```

process 1



4 KB



```
0x1010: movl 0x8(%rbp), %edi  
0x1013: addl $0x3, %edi  
0x1019: movl %edi, 0x8(%rbp)
```

8 KB

why didn't we have to
rewrite the stack addr?

process 2



12 KB



```
0x3010: movl 0x8(%rbp), %edi  
0x3013: addl $0x3, %edi  
0x3019: movl %edi, 0x8(%rbp)
```

16 KB

Problem: How to Run Multiple Processes?

Approaches (covered today):

Time Sharing

Static Relocation

Base

Base+Bounds

Segmentation

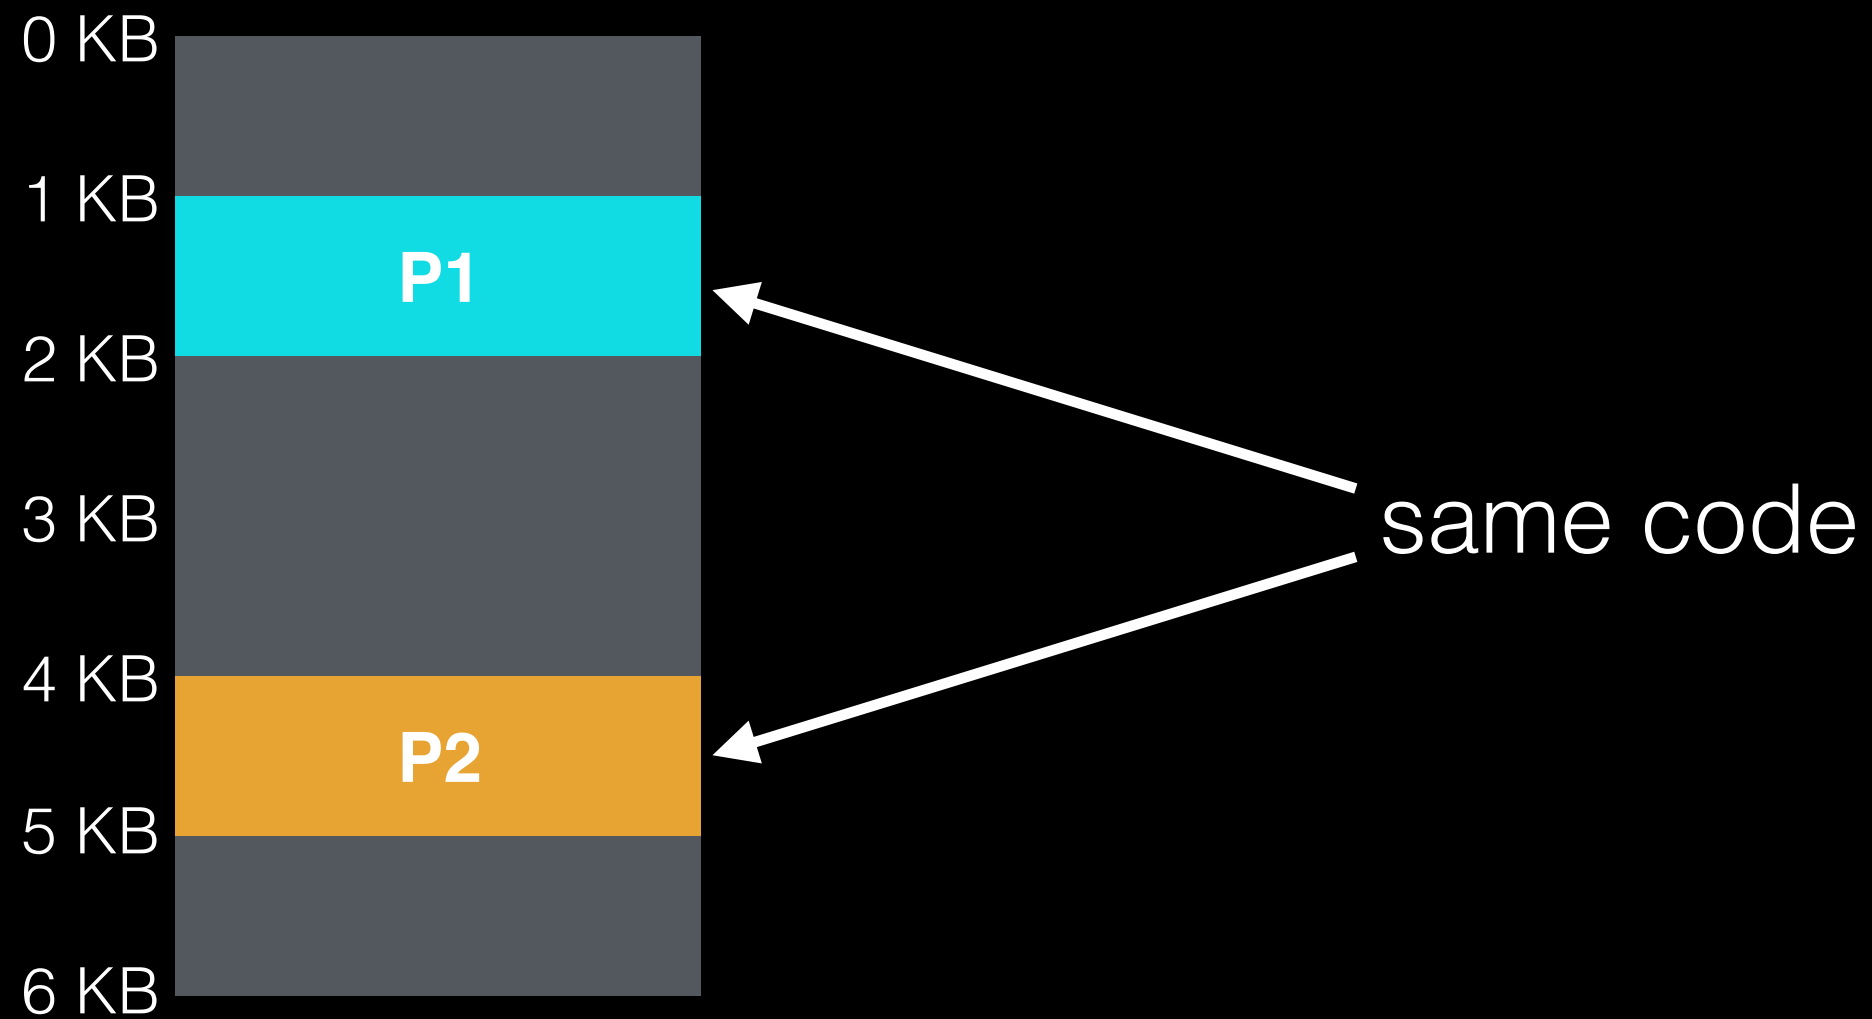
Base

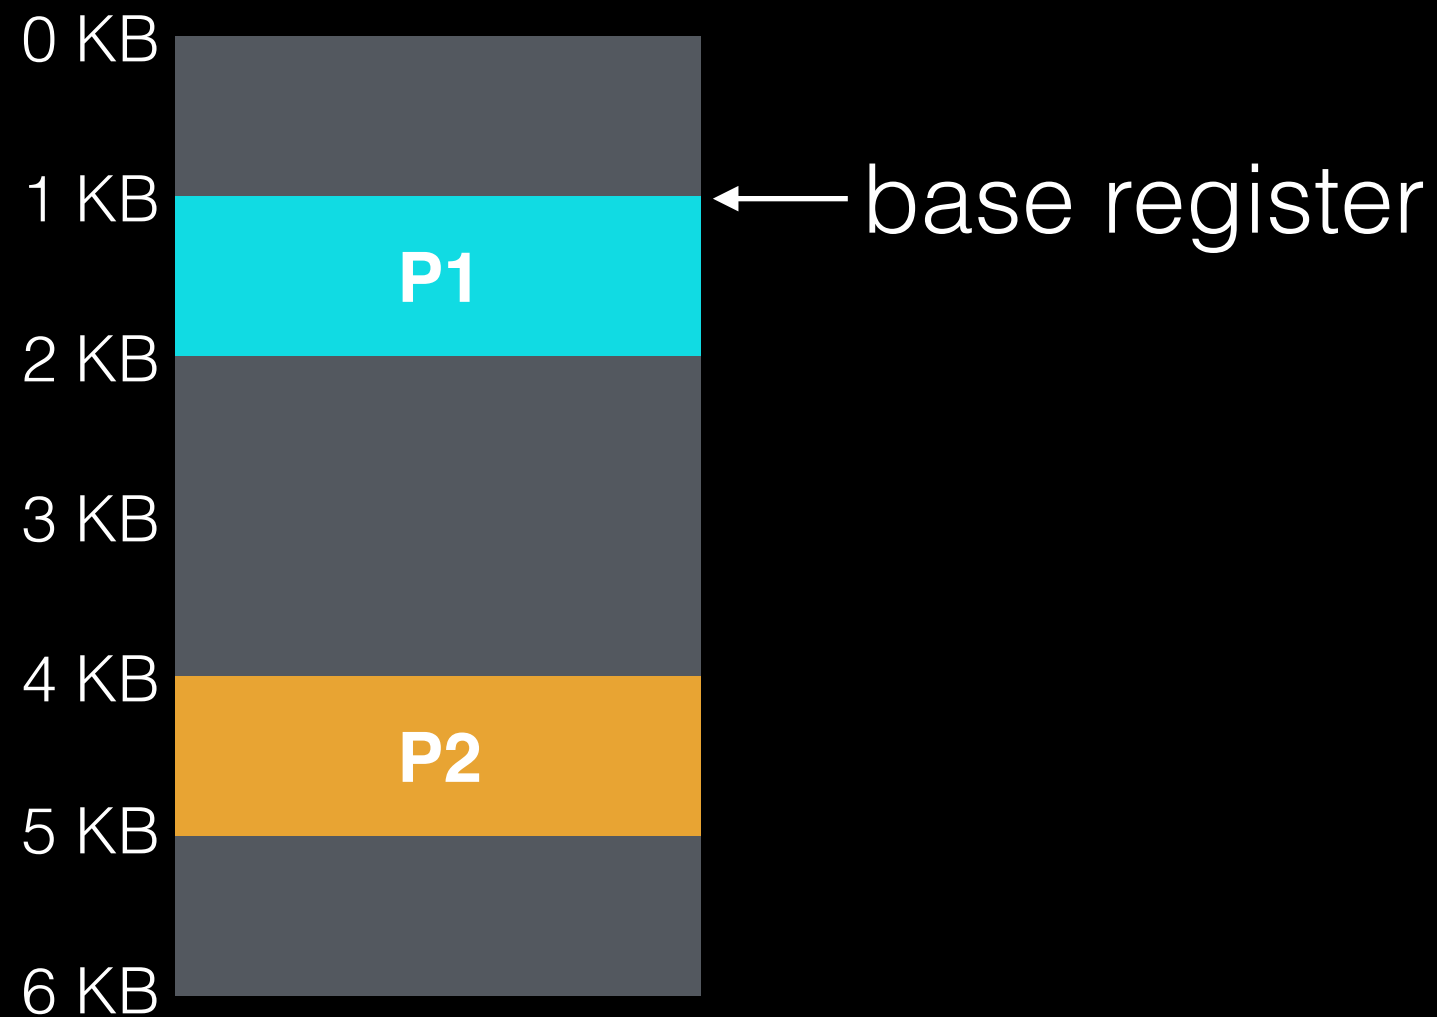
Idea: **translate** virtual addresses to physical by adding a fixed offset each time.

Store offset in a **base register**.


Each process has a different value in the base register when running.

This is a “**dynamic relocation**” technique





P1 is running



← base register

P2 is running




Virtual	Physical
P1: load 100, R1	



Virtual	Physical
P1: load 100, R1	load 1124, R1




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1


Who Controls the Base Register?

Who should **do translation** with base register?

(1) process, (2) OS, or (3) HW

Who should **modify** the base register?

(1) process, (2) OS, or (3) HW



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1

Can P2 hurt P1?
Can P1 hurt P2?



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	store 4096, R1

Can P2 hurt P1?
Can P1 hurt P2?

Problem: How to Run Multiple Processes?

Approaches (covered today):

Time Sharing

Static Relocation

Base

Base+Bounds

Segmentation

Base+Bounds

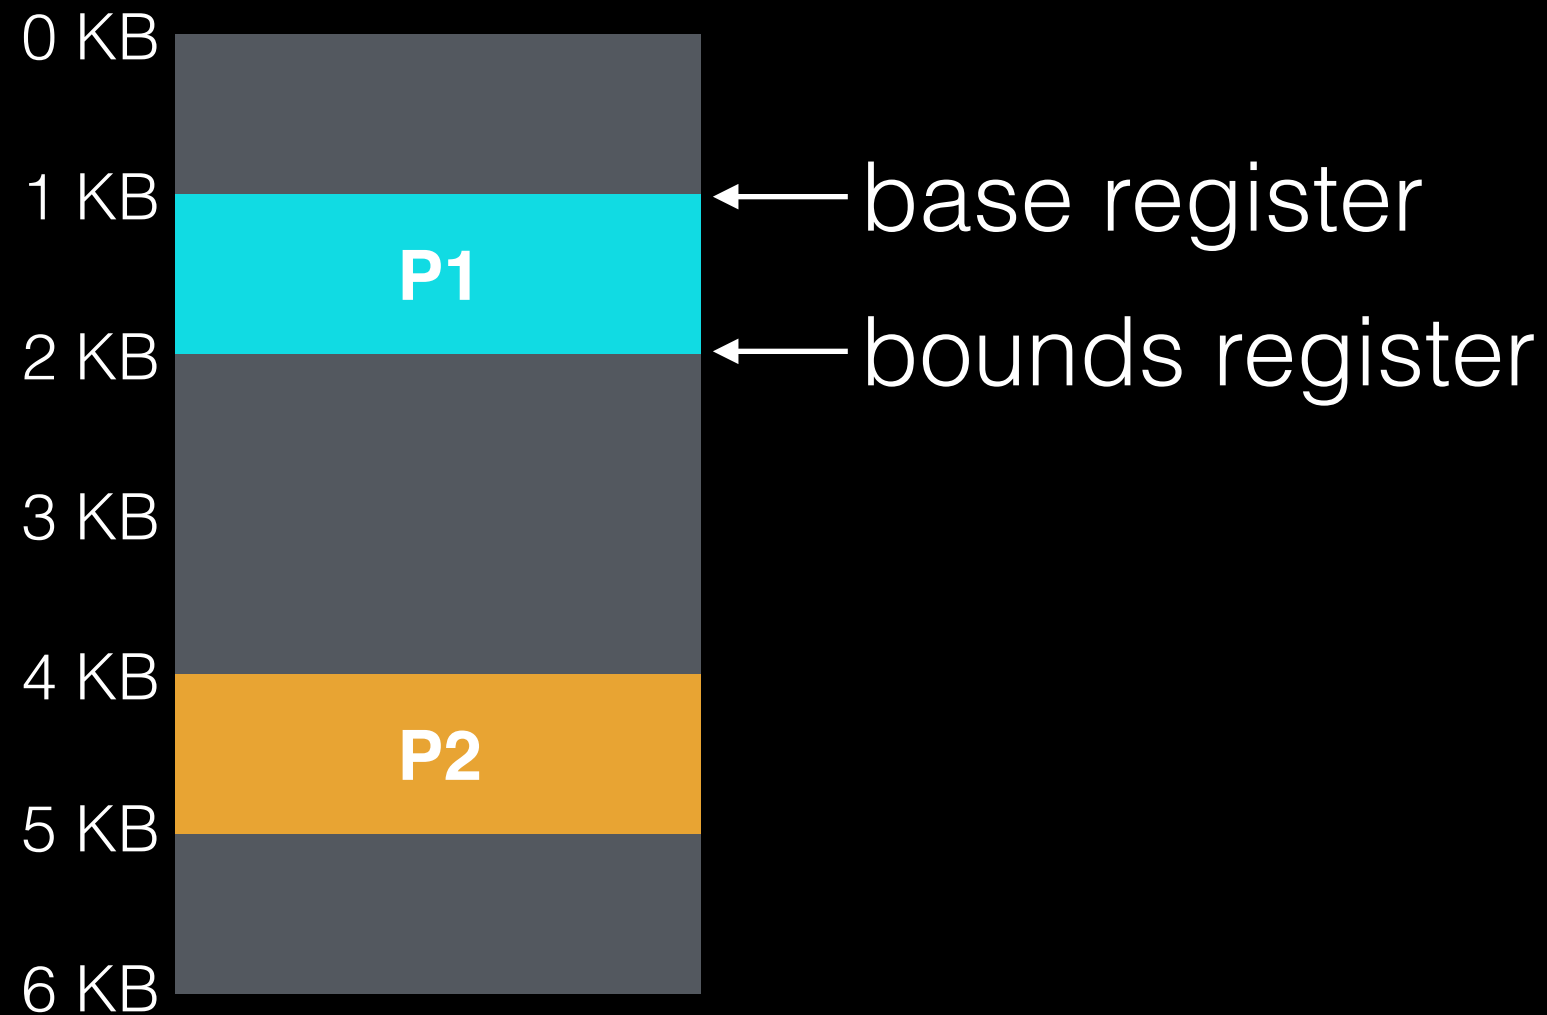
Idea: contain the address space with a bounds register marking the largest physical address

Base register: smallest physical addr


Bounds register: largest physical addr

OSTEP!

What happens if you load/store after bounds?




P1 is running




P2 is running

← base register
← bounds register




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	

Can P1 hurt P2?



Can P1 hurt P2?

Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	interrupt OS!



Can P1 hurt P2?

Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	interrupt OS!

Base+Bounds Pros/Cons

Pros?

Cons?

Base+Bounds Pros/Cons

Pros?

- fast + simple
- little bookkeeping overhead (2 registers / proc)

Cons?

- not flexible
- wastes memory for large address spaces

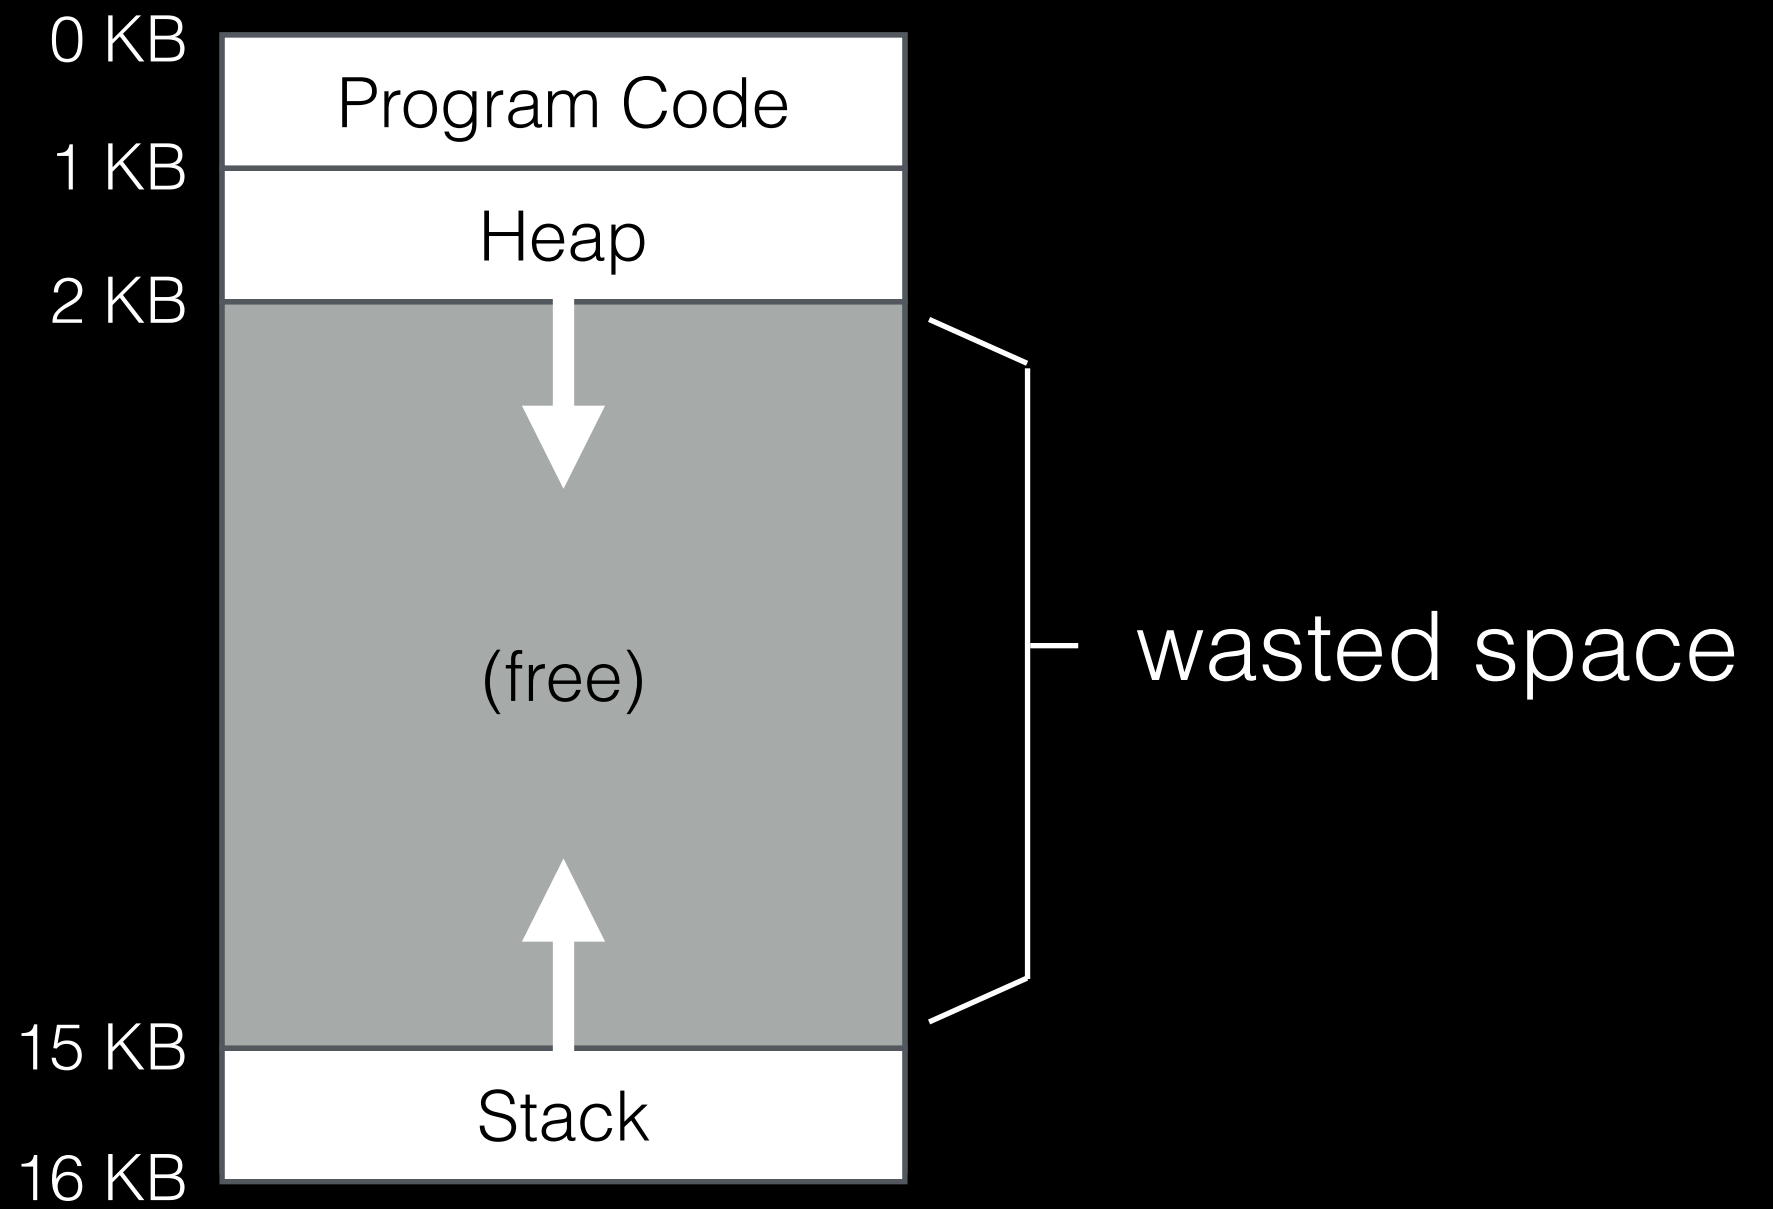
Base+Bounds Pros/Cons

Pros?

- fast + simple
- little bookkeeping overhead (2 registers / proc)

Cons?

- not flexible
- **wastes memory** for large address spaces



Problem: How to Run Multiple Processes?

Approaches (covered today):

Time Sharing

Static Relocation

Base

Base+Bounds

Segmentation

Segmentation

Idea: generalize base+bounds

Each base+bound pair is a *segment*

Use *different segments* for heap and memory

- how does this help?
- requires more registers!

Resize segments as needed


- how does this help?

Multi-segment translation


One (broken) approach:

- have **no gaps** in virtual addresses
- map as many low addresses to the first segment as possible, then as many as possible to the second (on so on)

A tricky example




A tricky example




Virtual	Physical
P1: load 100, R1	load 1124, R1

A tricky example




Virtual	Physical
P1: load 100, R1	load 1124, R1
P1: load 1024, R1	

A tricky example




Virtual	Physical
P1: load 100, R1	load 1124, R1
P1: load 1024, R1	load 4096, R1

A tricky example




Virtual	Physical
P1: load 100, R1	load 1124, R1
P1: load 1024, R1	load 4096, R1

A tricky example



Virtual	Physical
P1: load 100, R1	load 1124, R1
P1: load 1024, R1	load 4096, R1
P1: load 1024, R1	

A tricky example



Virtual	Physical
P1: load 100, R1	load 1124, R1
P1: load 1024, R1	load 4096, R1
P1: load 1024, R1	load 2048, R1

Multi-segment translation

One (correct) approach:

- break virtual addresses into two parts
- one part indicates segment
- one part indicates offset within segment

Virtual Address


For example, say addresses are 14 bits.
Use 2 bits for **segment**, 12 bits for **offset**

An address might look like 201E

Virtual Address

For example, say addresses are 14 bits.
Use 2 bits for **segment**, 12 bits for **offset**

An address might look like **2 01E**



segment 2 offset 30

Virtual Address

For example, say addresses are 14 bits.
Use 2 bits for **segment**, 12 bits for **offset**

An address might look like **2 01E**

Choose some segment numbering, such as

- 0: code+data

- 1: heap

- 2: stack

What is the segment/offset?

Segment numbers:

0: code+data

1: heap

2: stack

10 0000 0001 0001 (binary)

110A (hex)

4096 (decimal)




Virtual	Physical
load 0x2010, R1	



Virtual	Physical
load 0x2010, R1	4KB + 16




Virtual	Physical
load 0x2010, R1	4KB + 16
load 0x1010, R1	



Virtual	Physical
load 0x2010, R1	4KB + 16
load 0x1010, R1	1KB + 16



Virtual	Physical
load 0x2010, R1	4KB + 16
load 0x1010, R1	1KB + 16
load 0x1100, R1	



Virtual	Physical
load 0x2010, R1	4KB + 16
load 0x1010, R1	1KB + 16
load 0x1100, R1	1KB + 256



Virtual	Physical
load 0x2010, R1	4KB + 16
load 0x1010, R1	1KB + 16
load 0x1100, R1	1KB + 256
load 0x1400, R1	



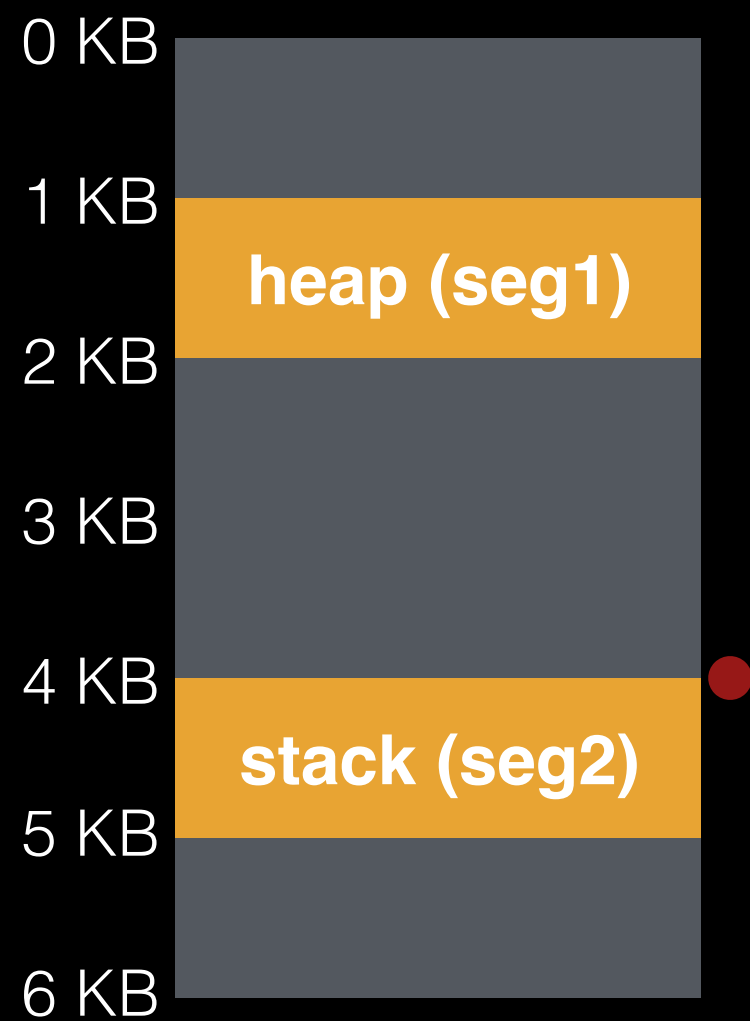
Virtual	Physical
load 0x2010, R1	4KB + 16
load 0x1010, R1	1KB + 16
load 0x1100, R1	1KB + 256
load 0x1400, R1	interrupt OS!

Stack Growth Problem

Example...



Virtual	Physical
load 0x2000, R1	



Virtual	Physical
load 0x2000, R1	4KB



Virtual	Physical
load 0x2000, R1	4KB



Virtual	Physical
load 0x2000, R1	4KB
load 0x2000, R1	



Virtual	Physical
load 0x2000, R1	4KB
load 0x2000, R1	3KB

Stack Growth Problem

Example...

Problem: $\text{phys} = \text{virt_offset} + \text{base_reg}$

phys is anchored to `base_reg`, which moves

Stack Growth Problem

Example...

Problem: $\text{phys} = \text{virt_offset} + \text{base_reg}$
phys is anchored to `base_reg`, which moves

Solution: anchor heap segment to `bounds_reg`:
 $\text{phys} = \text{bounds_reg} - (\text{max_offset} - \text{virt_offset})$

Stack Growth Problem

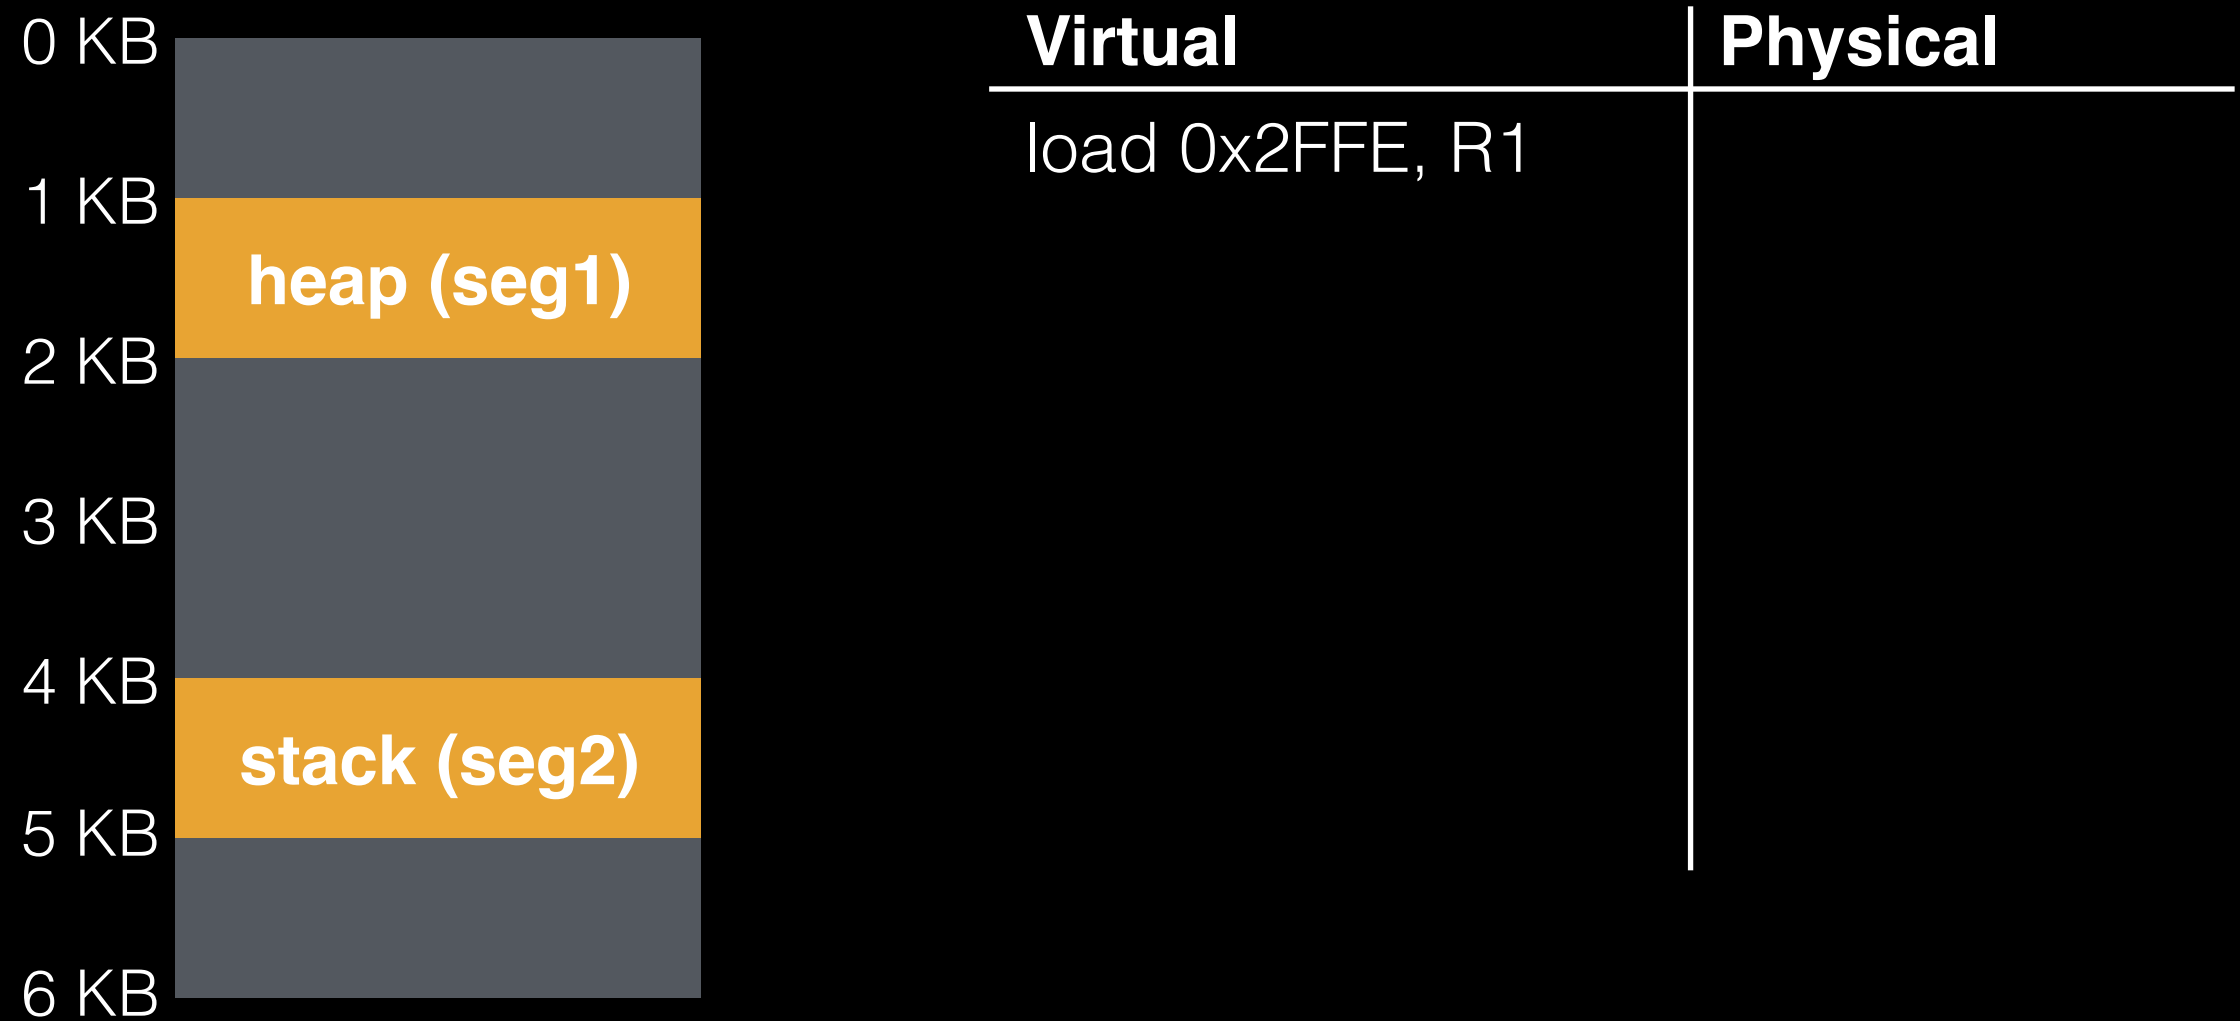
Example...

Problem: $\text{phys} = \text{virt_offset} + \text{base_reg}$
phys is anchored to `base_reg`, which moves

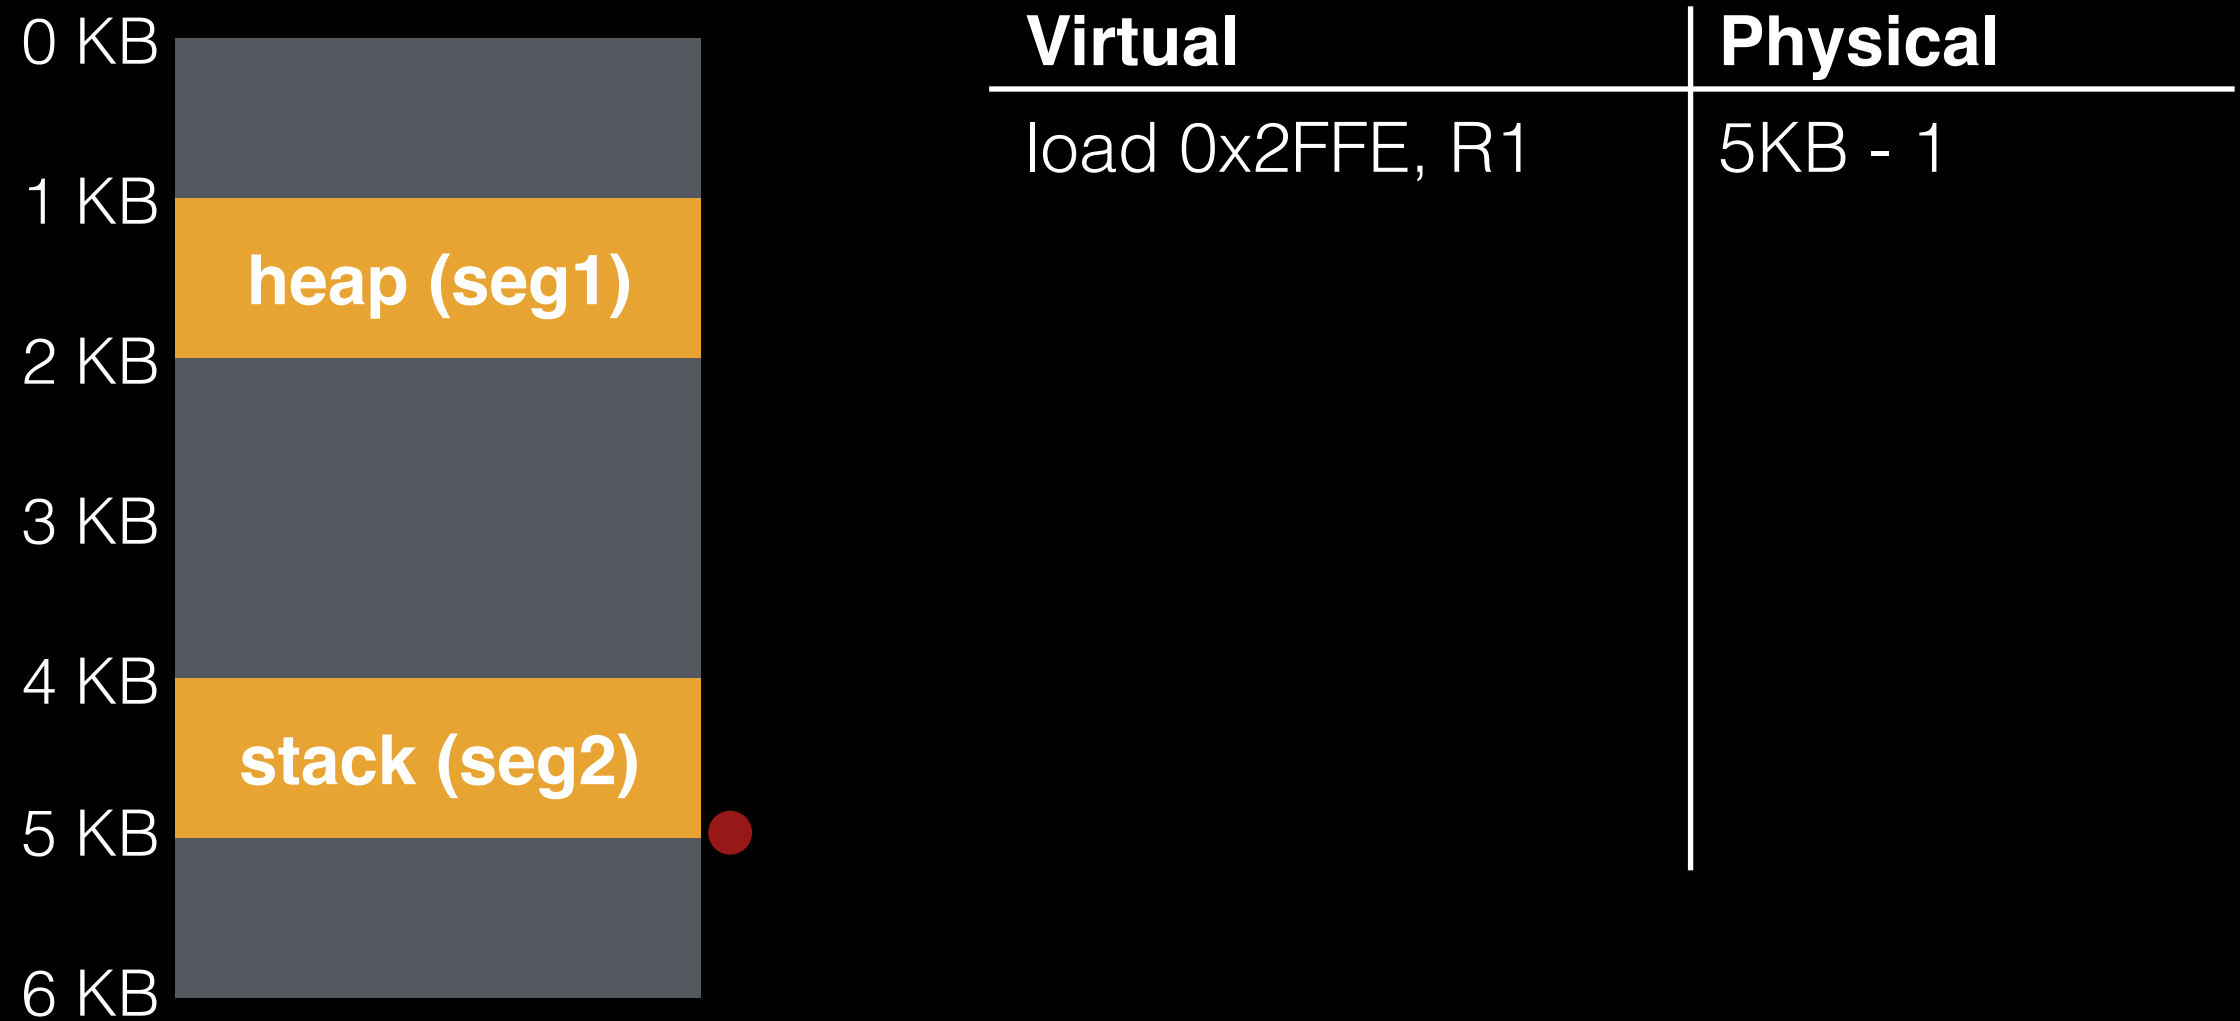
Solution: anchor heap segment to `bounds_reg`:
 $\text{phys} = \text{bounds_reg} - (\text{max_offset} - \text{virt_offset})$

Example (with `max_offset = FFF`)...

stack's max_offset = FFF



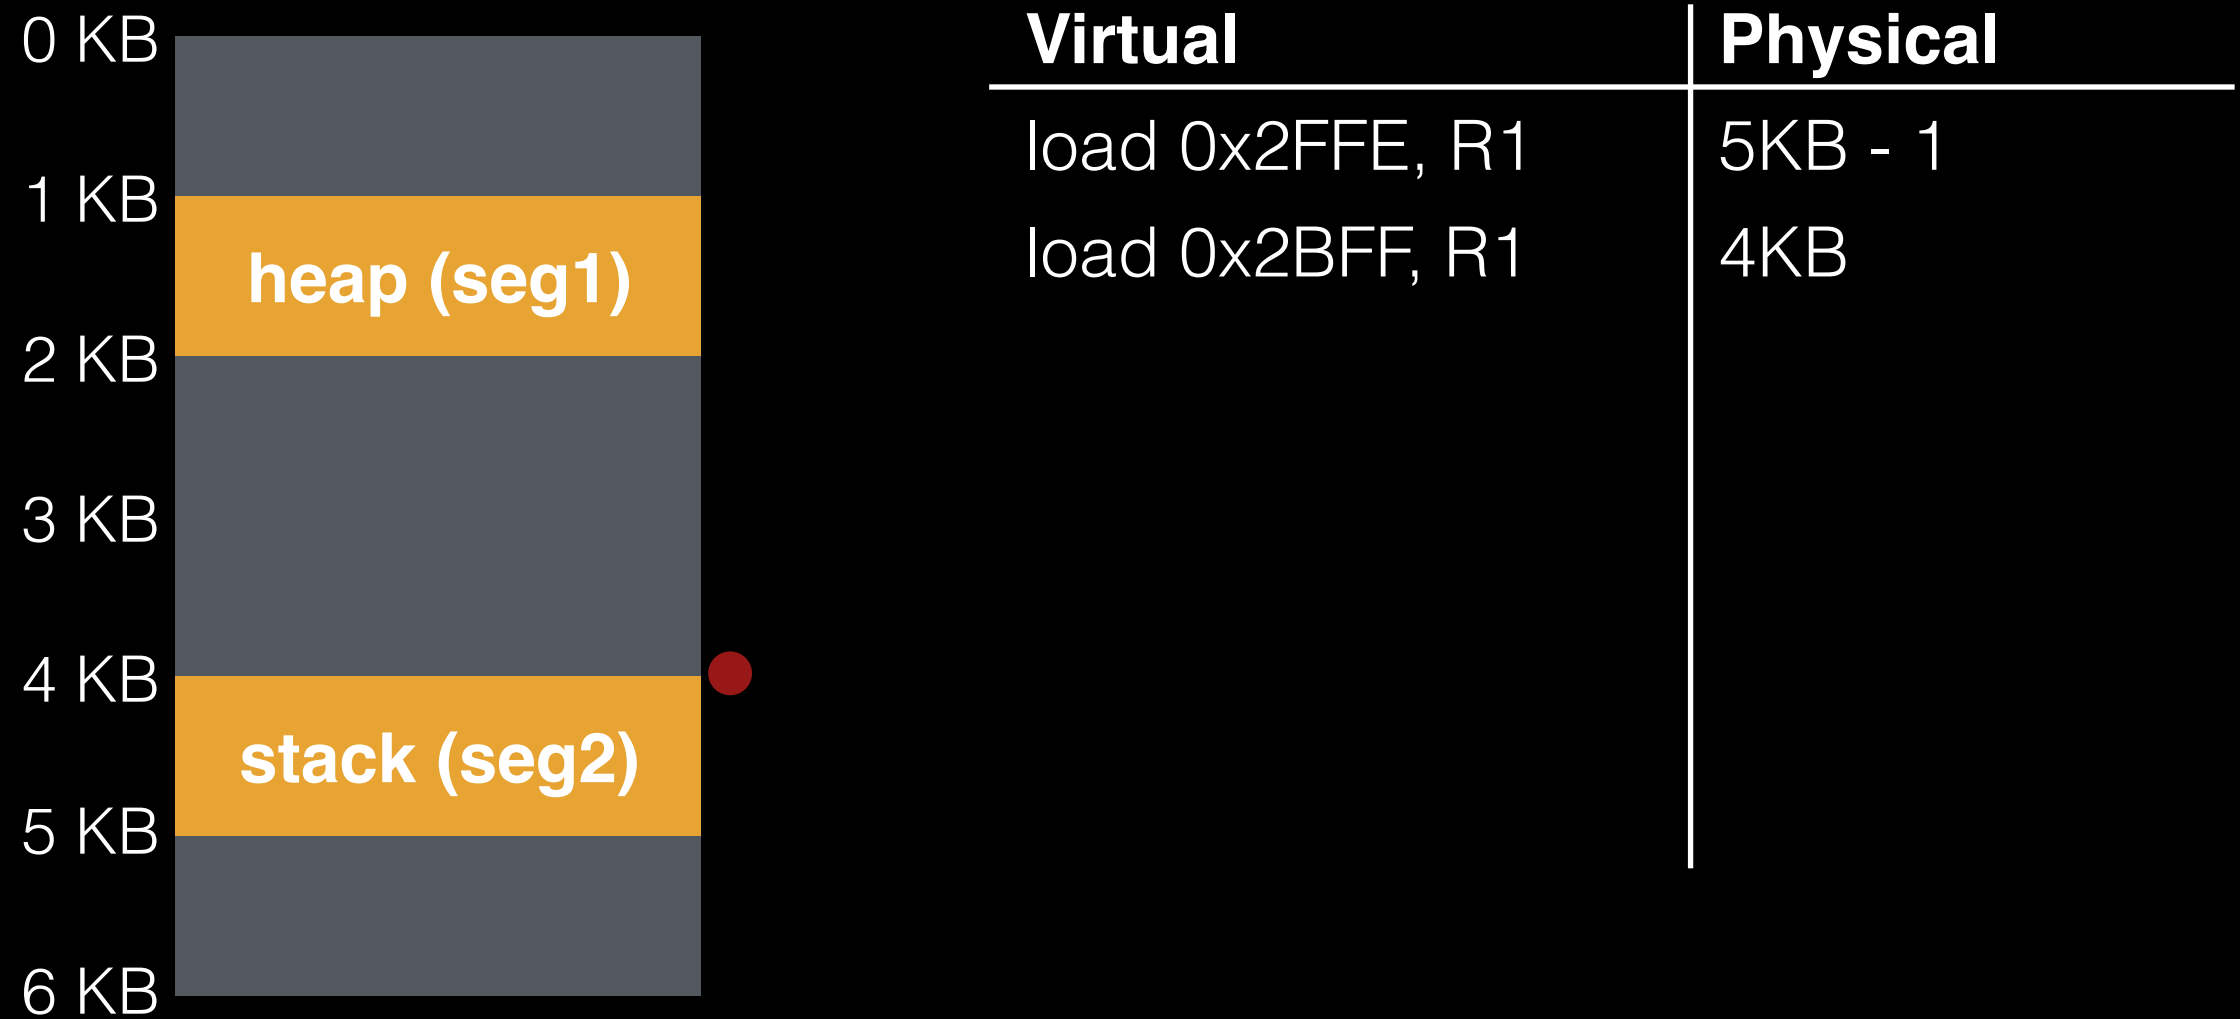
stack's max_offset = FFF



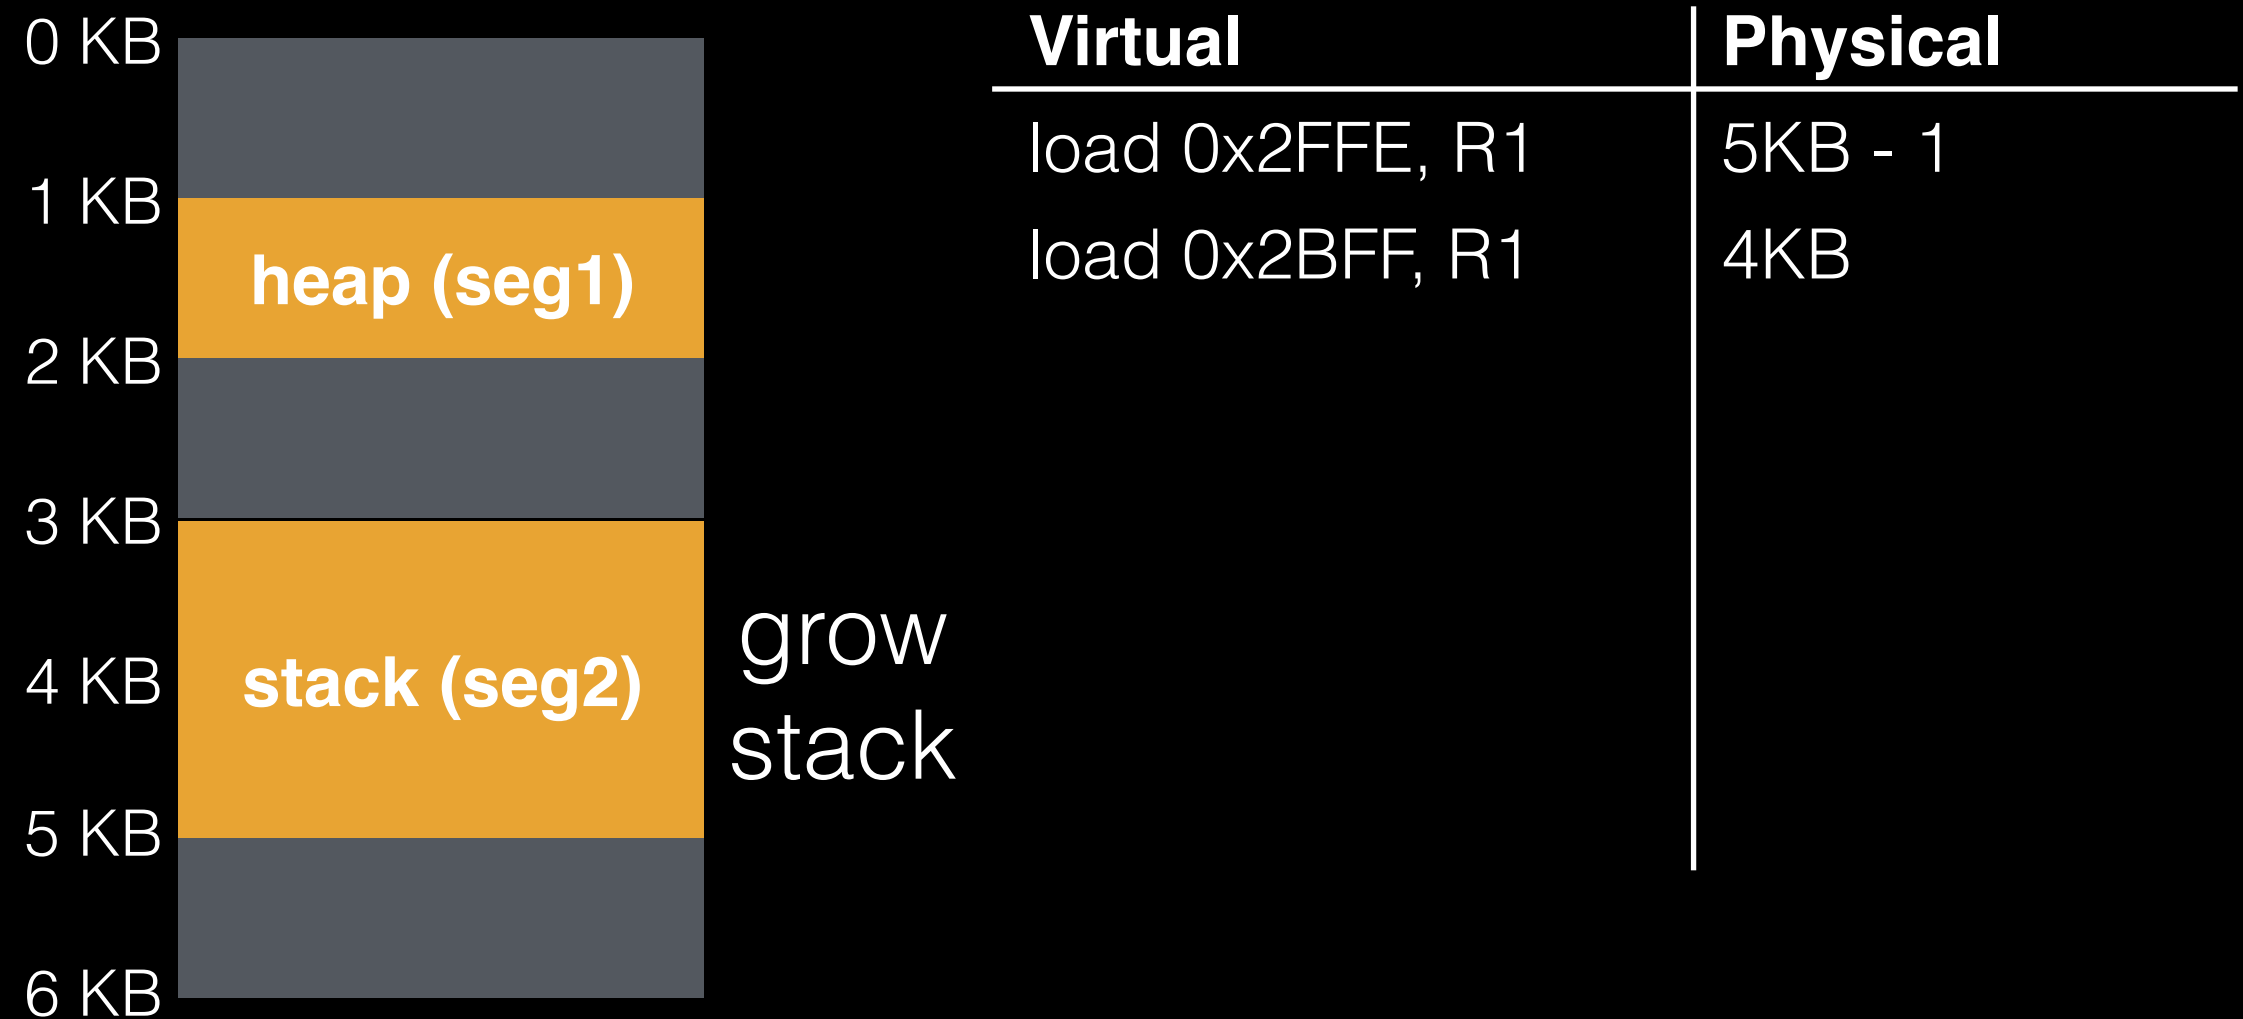
stack's max_offset = FFF

		Virtual	Physical
0 KB			
1 KB	heap (seg1)	load 0x2FFE, R1	5KB - 1
2 KB		load 0x2BFF, R1	
3 KB			
4 KB	stack (seg2)		
5 KB			
6 KB			

stack's max_offset = FFF



stack's max_offset = FFF



stack's max_offset = FFF

		Virtual	Physical
0 KB			
1 KB	heap (seg1)	load 0x2FFE, R1	5KB - 1
2 KB		load 0x2BFF, R1	4KB
3 KB		load 0x2BFF, R1	
4 KB	stack (seg2)		
5 KB			
6 KB			

stack's max_offset = FFF

		Virtual	Physical
0 KB			
1 KB	heap (seg1)	load 0x2FFE, R1	5KB - 1
2 KB		load 0x2BFF, R1	4KB
3 KB		load 0x2BFF, R1	4KB
4 KB	stack (seg2)		
5 KB			
6 KB			

Translation Summary

Heap: $\text{phys} = \text{base_reg} + \text{virt_offset}$

Stack: $\text{phys} = \text{bounds_reg} - (\text{max_offset} - \text{virt_offset})$


Anchors:

- for heap, anchor **smallest address** to **base register**
- for stack, anchor **biggest address** to **bounds register**

Code Sharing

Idea: make base/bounds for
the code of several processes
point to the same physical mem

Careful: need extra protection!



Segmentation Pros/Cons

Pros?

- supports sparse address space
- code sharing
- fine grained protection

Cons?

- external fragmentation

Conclusion

HW+OS work together to trick processes, giving the illusion of private memory

Adding CPU registers for base+bounds extends LDE, so translation is fast (does not always need OS)

Next time: solve fragmentation with paging
